

Introduction to C Programming

Lecture 14: basic pattern recognition II

Wenjin Wang
wangwj3@sustech.edu.cn

12-23-2022

Course syllabus

Nr.	Lecture	Date
1	Introduction	2022.9.9
2	Basics	2022.9.16
3	Decision and looping	2022.9.23
4	Array & string	2022.9.30
5	Functions	2022.10.9 (補)
6	Pointer	2022.10.14
7	Self-defined types	2022.10.21
8	I/O	2022.10.28

Nr.	Lecture	Date
9	Head files	2022.11.4

10	Review of lectures I	2022.11.25
11	Review of lectures II	2022.12.2
12	Review of lectures III	2022.12.9

13	AI in C programming	2022.12.16
14	AI in C programming	2022.12.23
15	Summary	2022.12.30

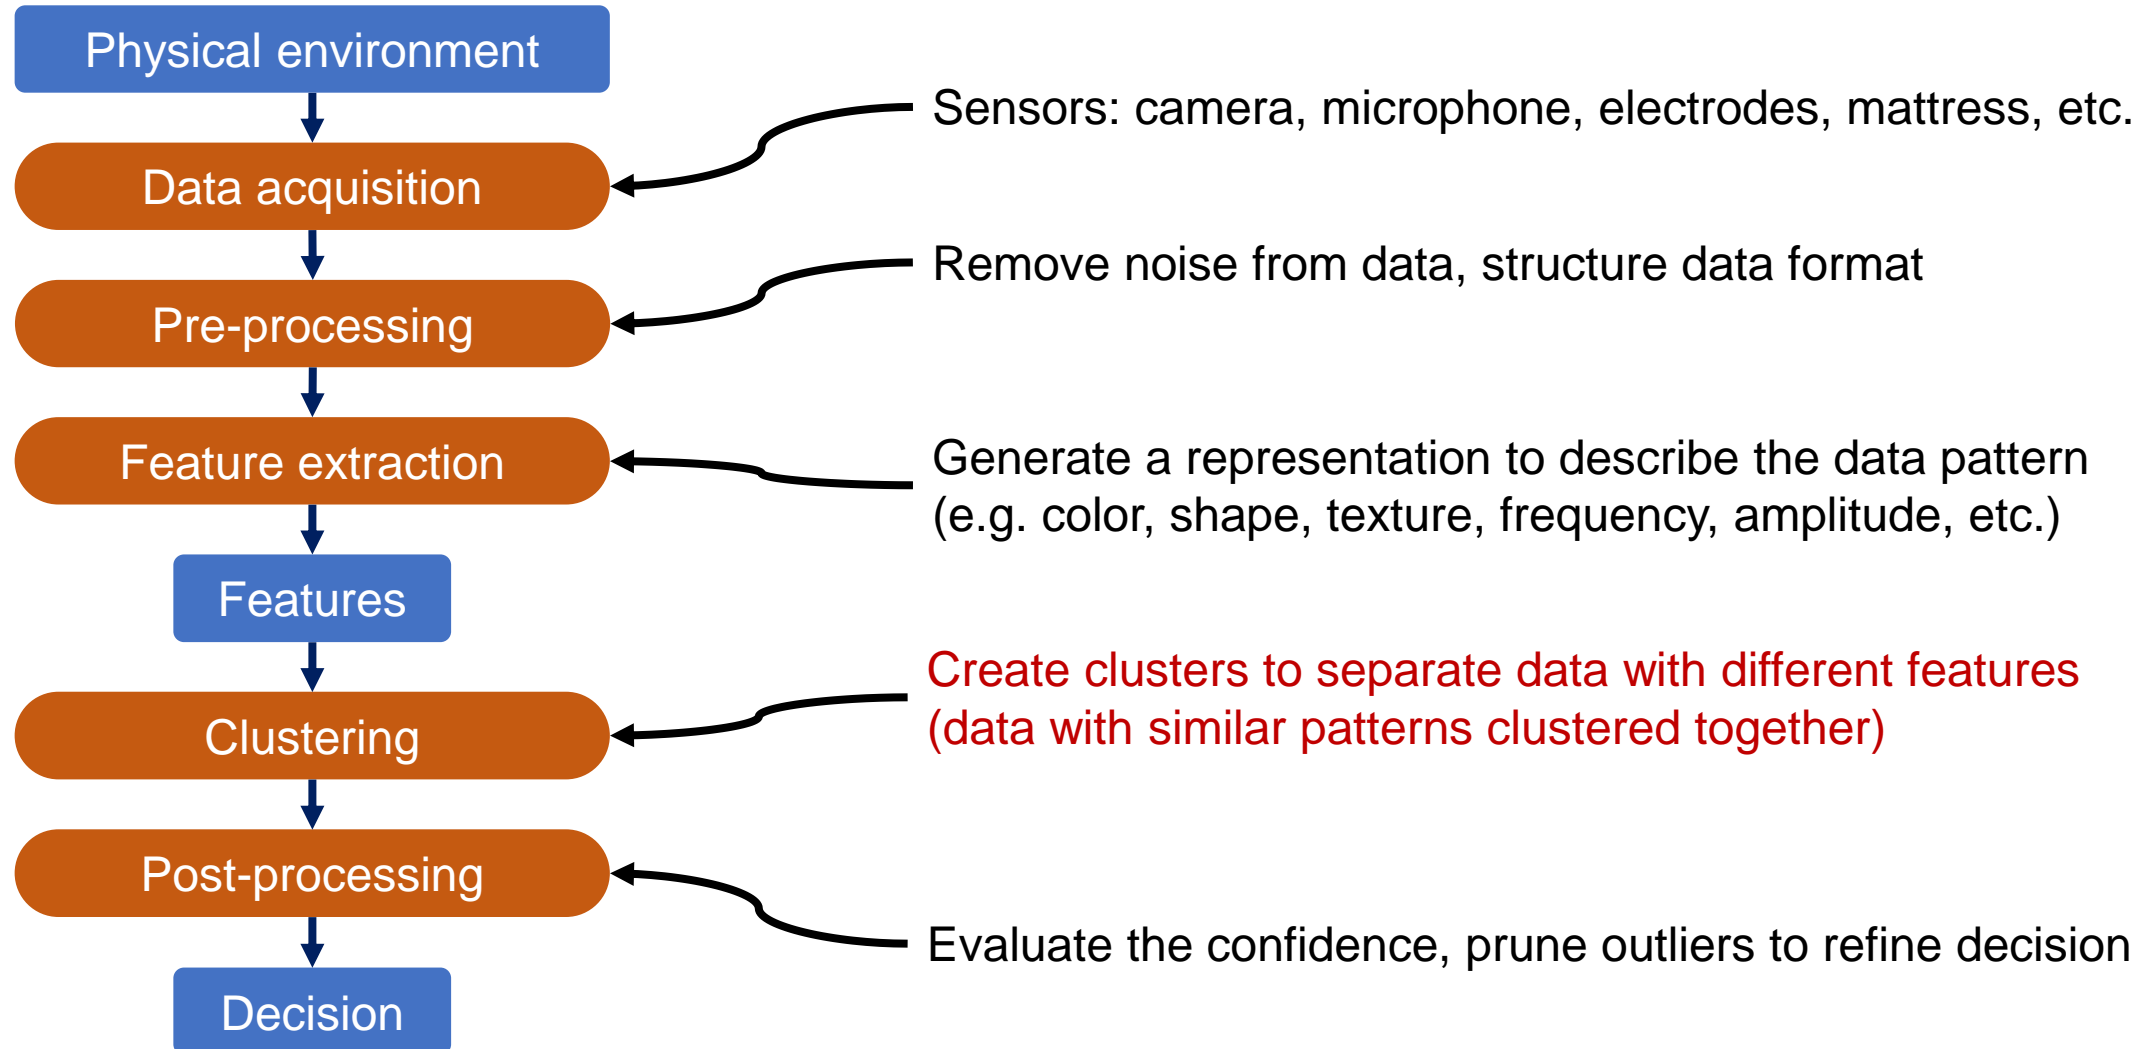
Objective of this lecture

You can play with basic AI algorithms!

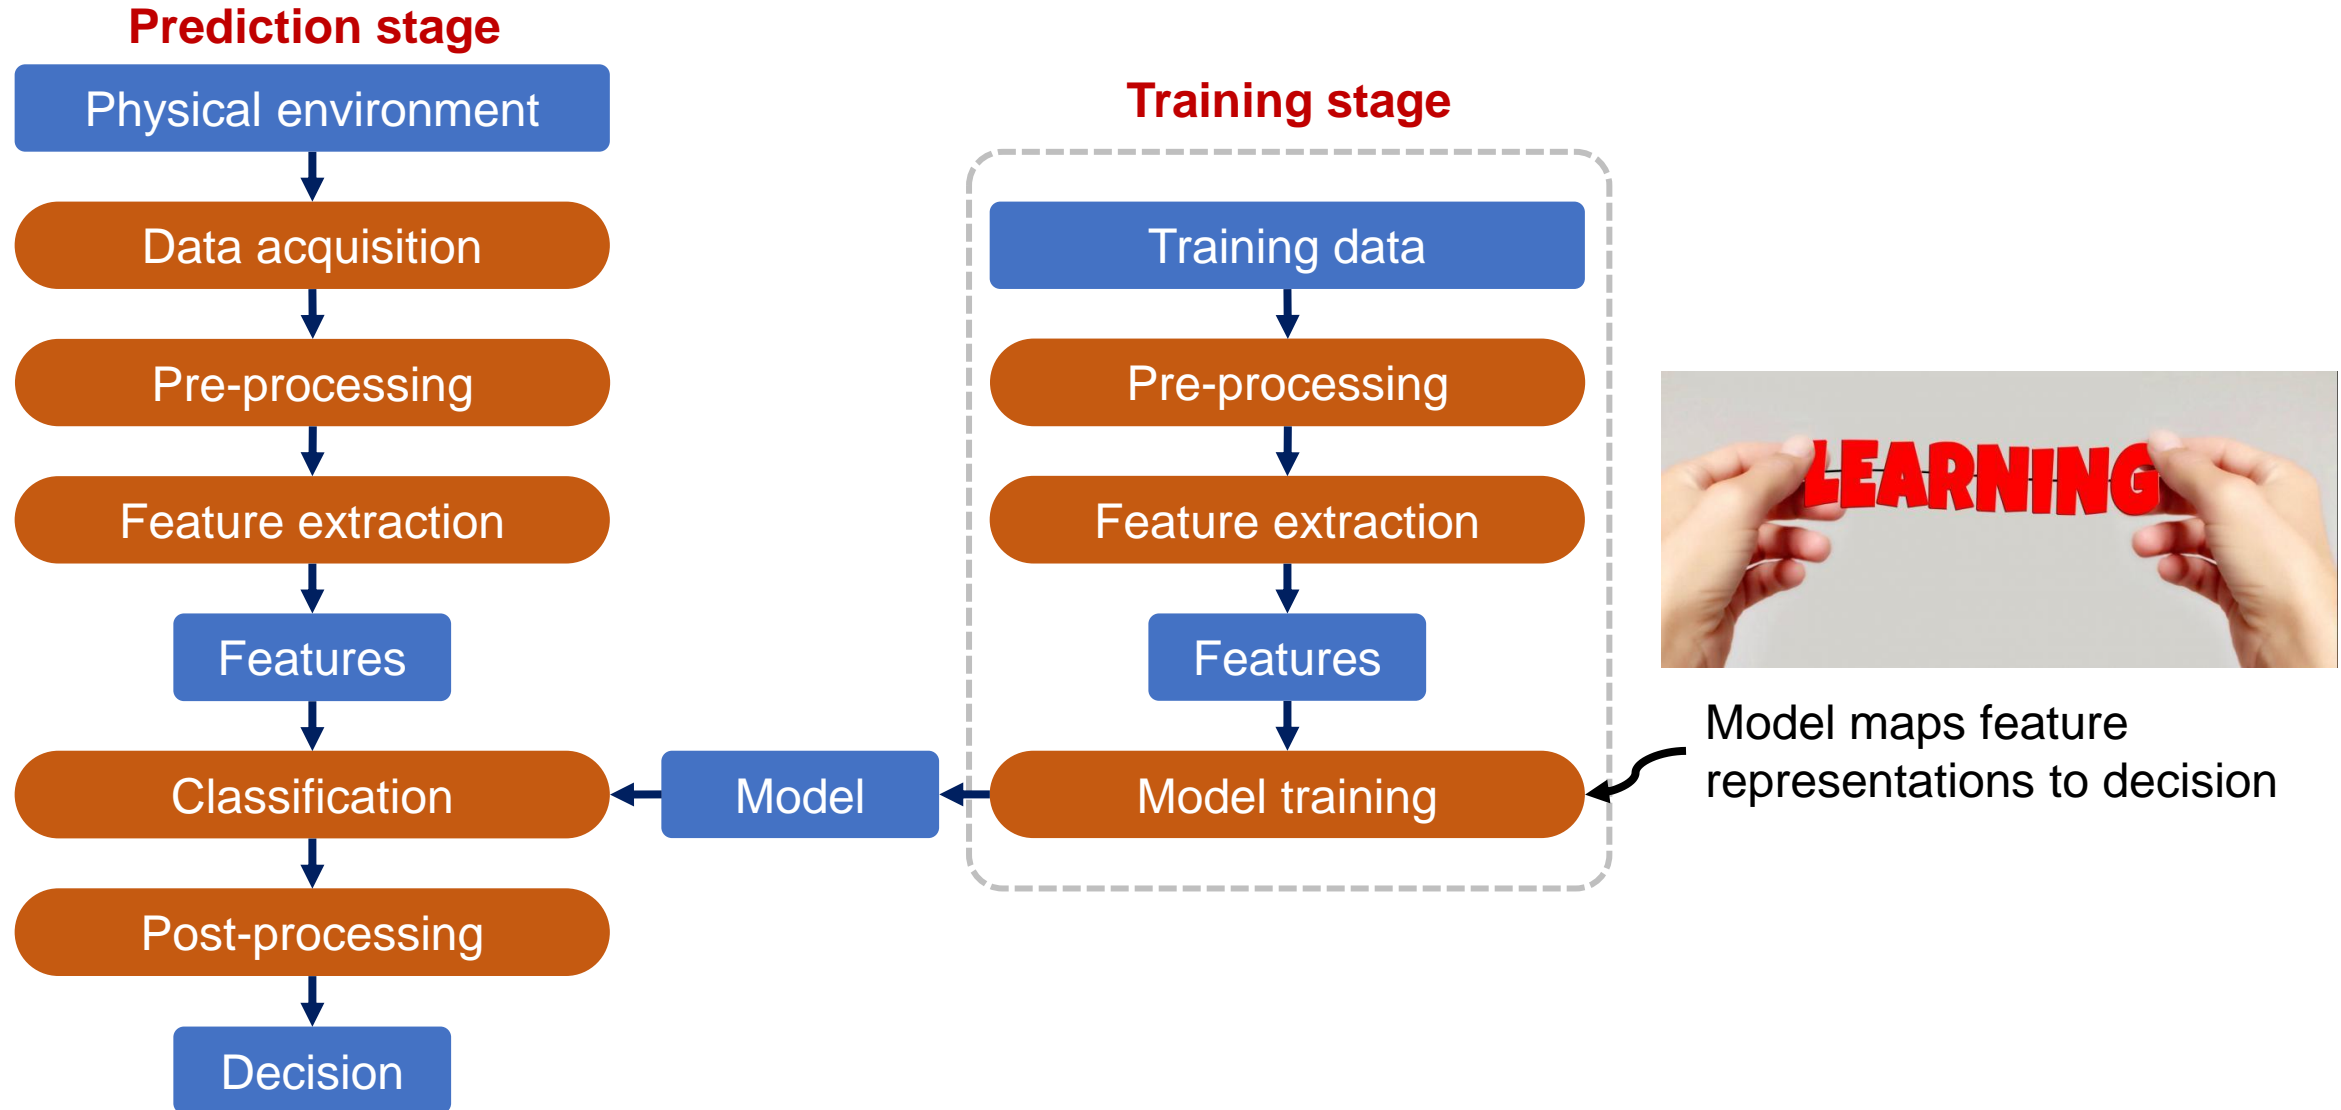
Recap last lecture

- **Pattern recognition** is the key to make machine intelligent
- Two major types of pattern recognition algorithms: **clustering** and **classification**
- Two representative clustering algorithms are given: **Kmeans** and **DBscan**
- **Kmeans** needs to manually define number of clusters (the value of K)
- **DBscan** needs to input scanning radius and number of neighbours
- You know how to install **OpenCV** to visualize the Kmeans learning process

Unsupervised clustering (聚类)



Supervised classification (分类)



Content

- 1. Basic classification (KNN & perceptron)**
- 2. Basics of CNN (deep learning!)**
- 3. Applications of CNN**

Content


- 1. Basic classification (KNN & perceptron)**
2. Basics of CNN (deep learning!)
3. Applications of CNN

What is KNN?

K Nearest Neighbor



K个



最邻近点

What is KNN?


K Nearest Neighbor



Majority of votes from
neighborhood

What is KNN?


In our fish sorting case, seabass and salmon are collected as training samples



What is KNN?

Seabass class: -1

Salmon class: 1




Define $K = 4$,
in 4 nearest neighbors, 3 are seabass, 1 is salmon

What is KNN?

Seabass class: -1

Salmon class: 1

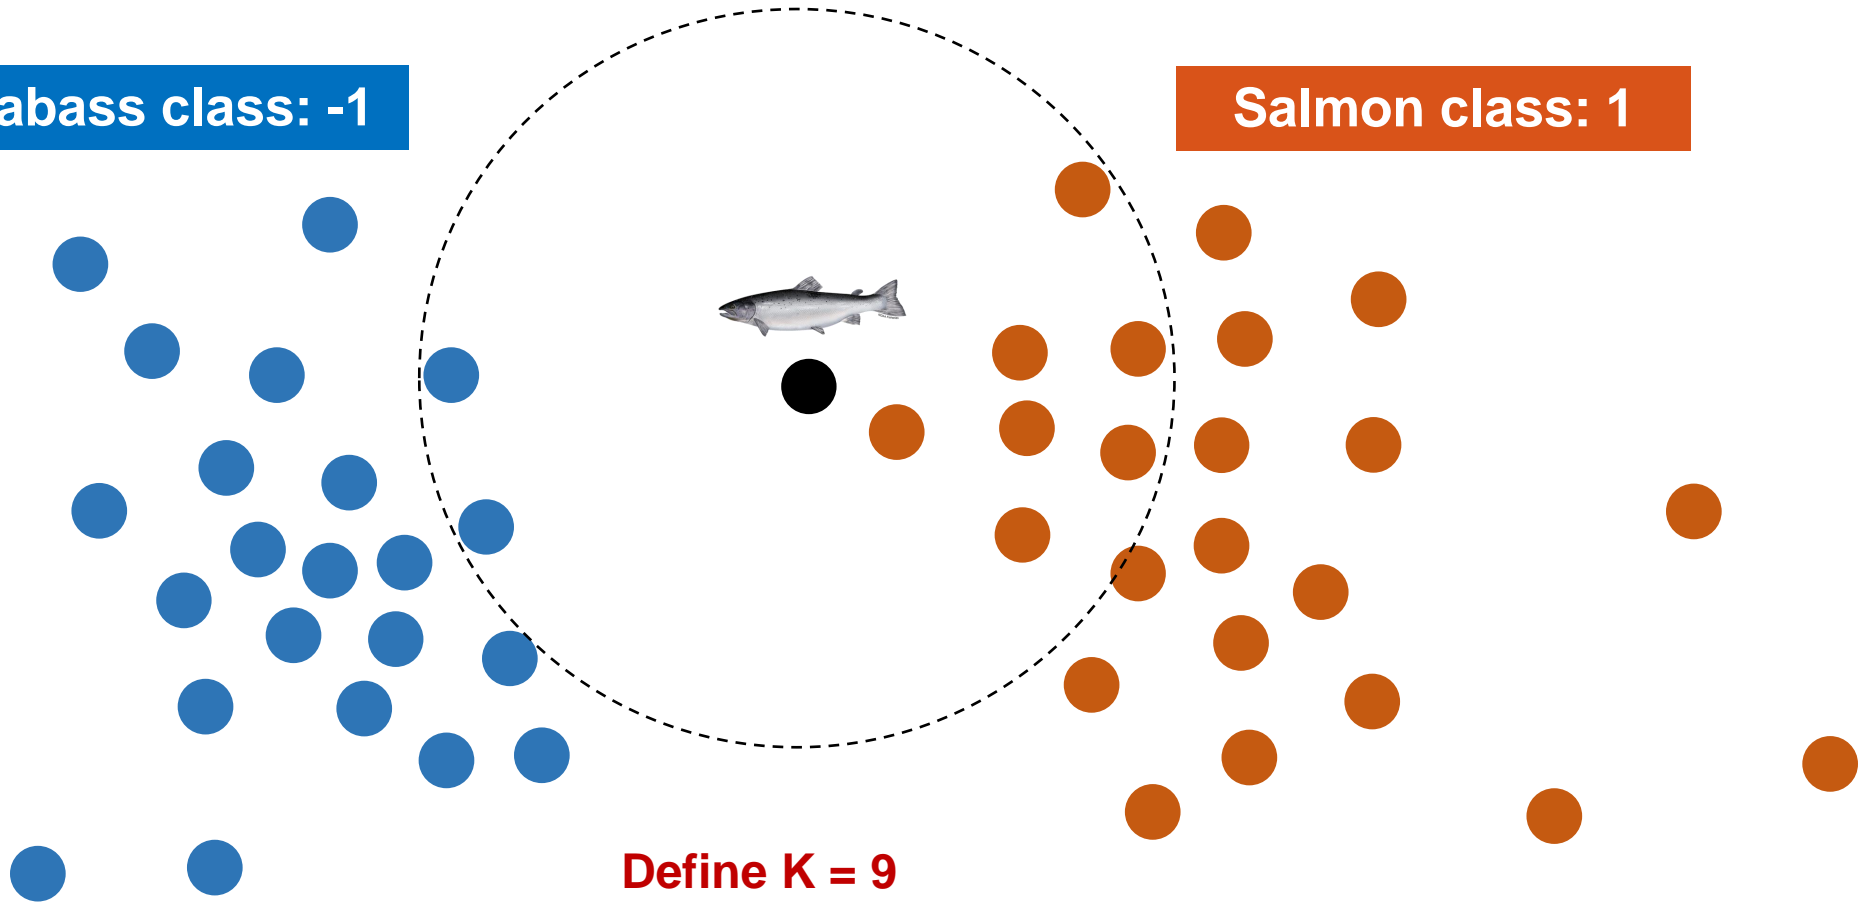


Define $K = 9$,
in 9 nearest neighbors, 7 are seabass, 2 are salmon

What is KNN?

Seabass class: -1

Salmon class: 1



Define $K = 9$
in 9 nearest neighbors, 2 are seabass, 7 are salmon

KNN definition

KNN finds a number of K nearest neighbors in the space and use the majority of class to label the sample.

Input: samples


- Calculate the distance between sample and training samples
- Sort the distances and select the K nearest samples
- Use the majority class of K nearest samples to label the sample.

Output: class of sample

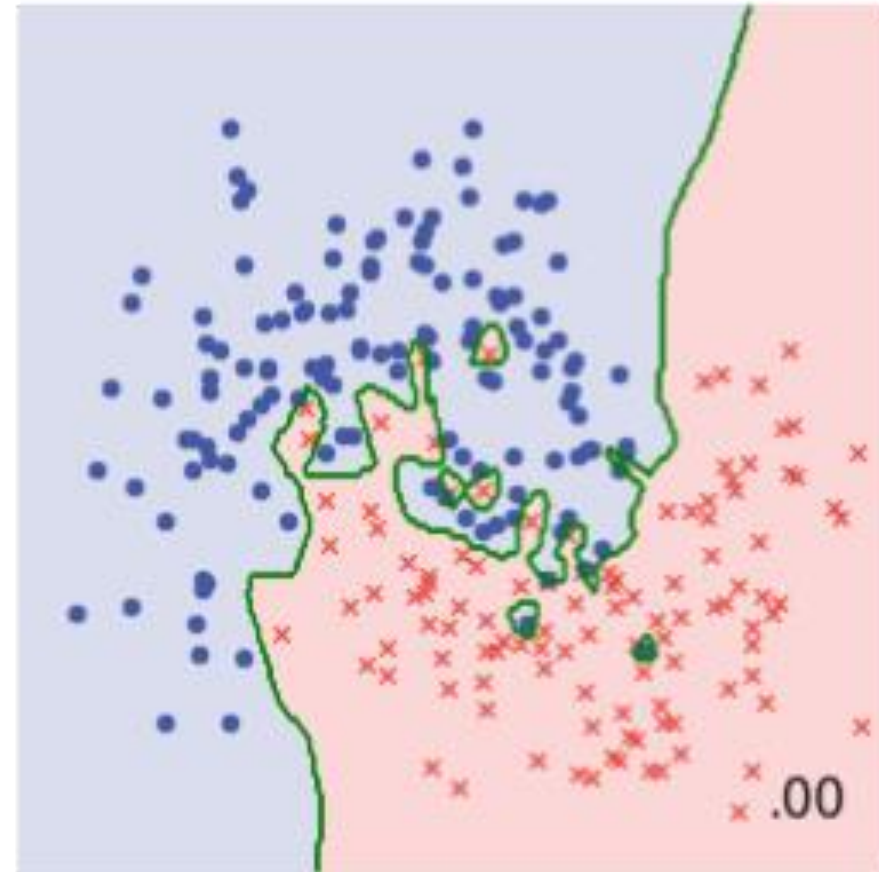
Effect of K

K is the ONLY parameter you need to define

k=99



k=1



Pros and cons of KNN

Pros

- Learning and implementation is extremely simple and intuitive
- Flexible decision boundaries (can be highly non-linear)

Cons

- Irrelevant or redundant features have negative impact
- Rely on the distance metric, cannot handle clusters with different densities

Applications of KNN


My MSc assignment 10 years ago at University of Amsterdam!

Develop an APP that search the information of an image taken by a phone!

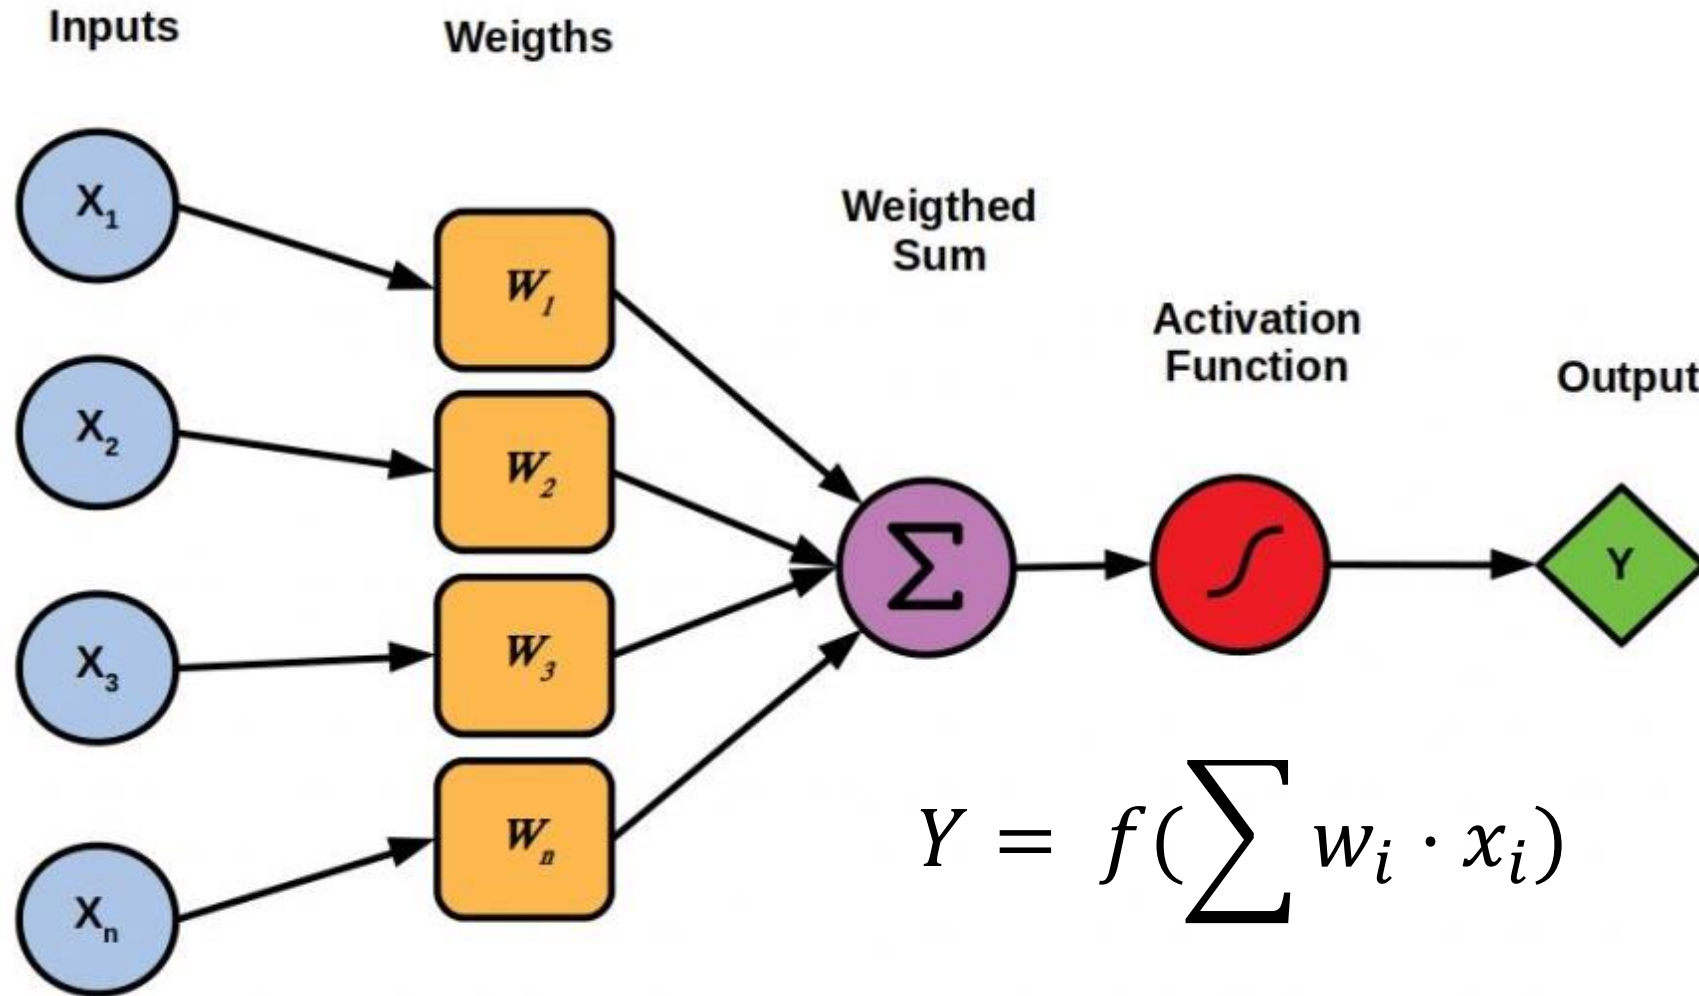
KNN algorithm was implemented to search the image.



Applications of KNN




What is perceptron?




Problem definition in perceptron

Seabass class: -1




Salmon class: 1



$x = (\text{colour}, \text{length});$
 $y = \text{label}; // -1 \text{ or } 1$

$$w \cdot x + b = 0$$

Problem definition in perceptron



Perceptron

Label:

$$y = \{-1, 1\}$$



Data:

$$x = (\text{color}, \text{length})$$



$$f(x) = \text{sign}(w \cdot x + b)$$



Model coefficients

Perceptron

Label:

$$y = \{-1, 1\}$$



Data:


$$x = (\text{color}, \text{length})$$



$$f(x) = \text{sign}(w \cdot x + b)$$




$$\text{sign}(x) = \begin{cases} +1, & x \geq 0 \\ -1, & x < 0 \end{cases}$$



Perceptron

$$f(x) = \text{sign}(w \cdot x + b)$$

$$w_1 \cdot x_1 + w_2 \cdot x_2 + \cdots + w_n \cdot x_n + b \quad \begin{cases} +1 \\ -1 \end{cases}$$

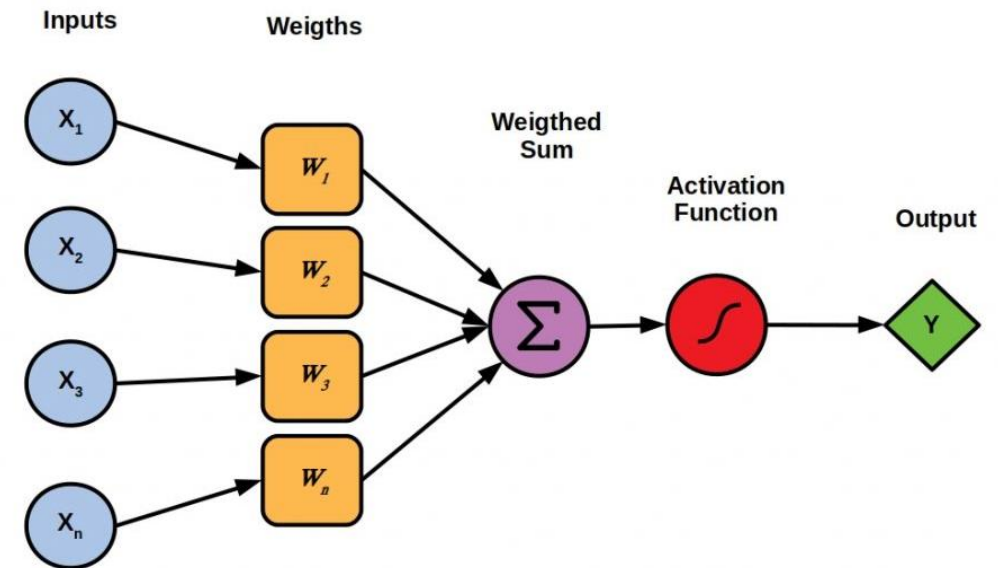

The diagram consists of two arrows. One arrow starts from the w_1 term in the equation above and points diagonally down and to the right towards the word "weights". The other arrow starts from the b term in the equation above and points diagonally down and to the left towards the word "bias".

weights

bias

Perceptron

$$f(x) = \text{sign}(w \cdot x + b)$$



Perceptron

$$f(x) = \text{sign}(w \cdot x + b)$$

Diagram illustrating the Perceptron function:

- $f(x)$ is labeled "Decision" with a blue arrow pointing to it.
- w and b are highlighted in red.
- x is labeled "From sensor" with a blue arrow pointing to it.
- A red bracket connects w and b , pointing down towards the text "How to get model parameters???"

How to get model parameters???

Train it!!!


Perceptron

How to train a perceptron?

**Minimize the error in
classification!**


Perceptron

Minimize the error in classification



Perceptron


Minimize the error in classification



$$\text{distance} = \frac{|a \cdot x + b \cdot y + c|}{\sqrt{a^2 + b^2}}$$

Perceptron

Minimize the error in classification




Using the notions in our task:

$$\text{distance} = \frac{|w_1 \cdot x_1 + w_2 \cdot x_2 + b|}{\sqrt{w_1^2 + w_2^2}}$$

Perceptron

Minimize the error in classification



Using the notions of our task:


$$\text{distance} = \frac{|w \cdot x + b|}{||w||}$$

$$= \frac{-y \cdot (w \cdot x + b)}{||w||}$$

↙ L2 norm

Perceptron

Minimize the error in classification




Sum of distances of all errors

$$\text{distance} = \sum \frac{-y \cdot (w \cdot x + b)}{||w||}$$

Norm is positive,
no impact on
results

Perceptron

Minimize the error in classification



Loss function can be defined as:

$$loss = - \sum y \cdot (w \cdot x + b)$$

Perceptron

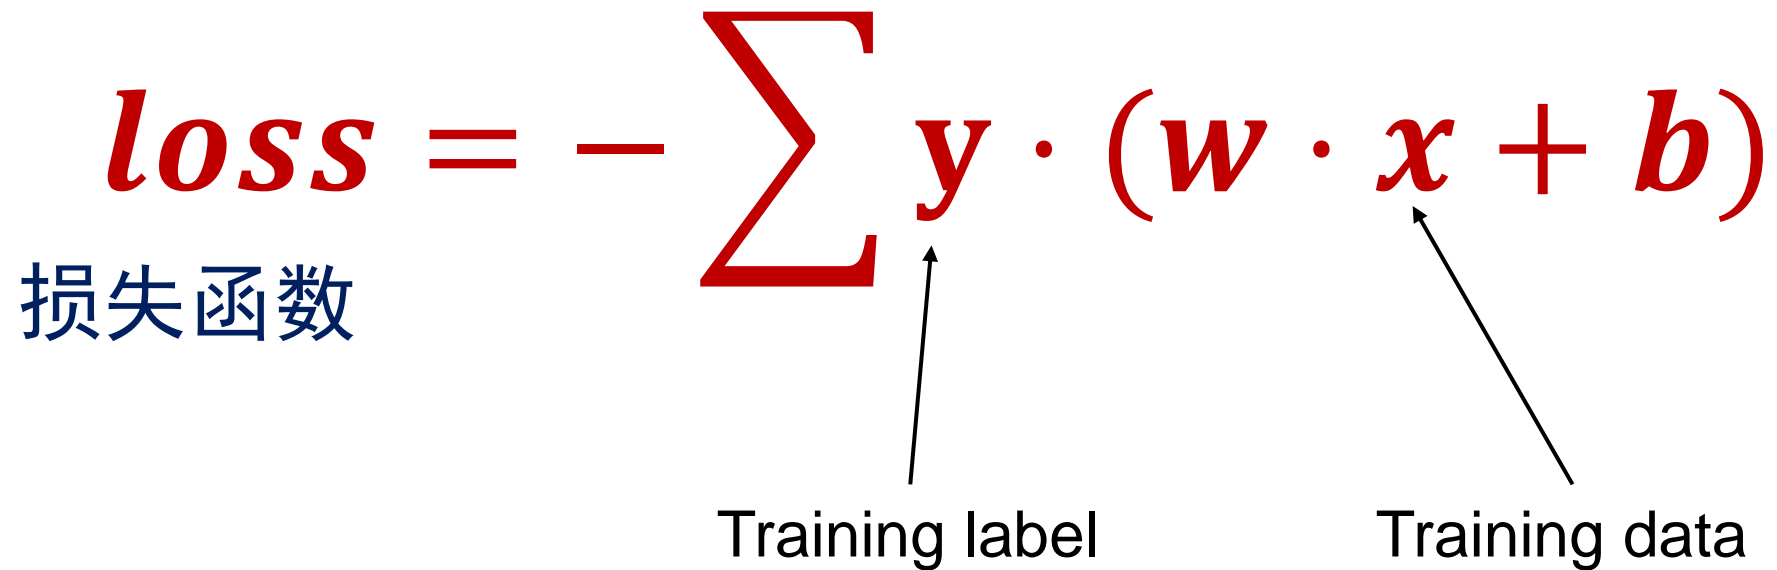
The goal is to find **w and b** that minimize loss

$$\text{loss} = - \sum \mathbf{y} \cdot (\mathbf{w} \cdot \mathbf{x} + b)$$


损失函数

Training label

Training data


The diagram shows the loss function formula in red: $\text{loss} = - \sum \mathbf{y} \cdot (\mathbf{w} \cdot \mathbf{x} + b)$. Below the formula, the Chinese text "损失函数" (Loss Function) is written in blue. Two arrows point from labels below to variables in the formula: one from "Training label" to the variable \mathbf{y} , and another from "Training data" to the variable \mathbf{x} .

Gradient descent optimization



Gradient descent optimization

Update the coefficients along the gradient direction of loss function is the fastest path to achieve the minimum!!!



Perceptron


The gradient of loss is calculated by
(partial) derivative

$$loss = - \sum y \cdot (w \cdot x + b) \quad \left\{ \begin{array}{l} \nabla_w L(w, b) = - \sum y_i \cdot x_i \\ \nabla_b L(w, b) = - \sum y_i \end{array} \right.$$

Perceptron

Update the coefficients step-wise along the gradient direction

Gradient descent



$$w \leftarrow w + \alpha \cdot \nabla_w L(w, b)$$

$$b \leftarrow b + \alpha \cdot \nabla_b L(w, b)$$

α ($0 < \alpha < 1$) learning rate

Pros and cons of perceptron

Pros:

- Simplicity
- Generally applicable for various tasks

Cons:

- Data must be linearly separable
- No guarantee on non-convex problem

Content

1. Basic classification (KNN & perceptron)
- 2. Basics of CNN (deep learning!)**
3. Applications of CNN

A basic introduction of CNN


CNN basics

(start with a concrete example)





CNN applications

(general ideas in image and video processing)



Assuming the case of facial mask check

- You must wear a **face mask** to enter the school!!!



Goal

- Camera-based classification of faces with/without mask

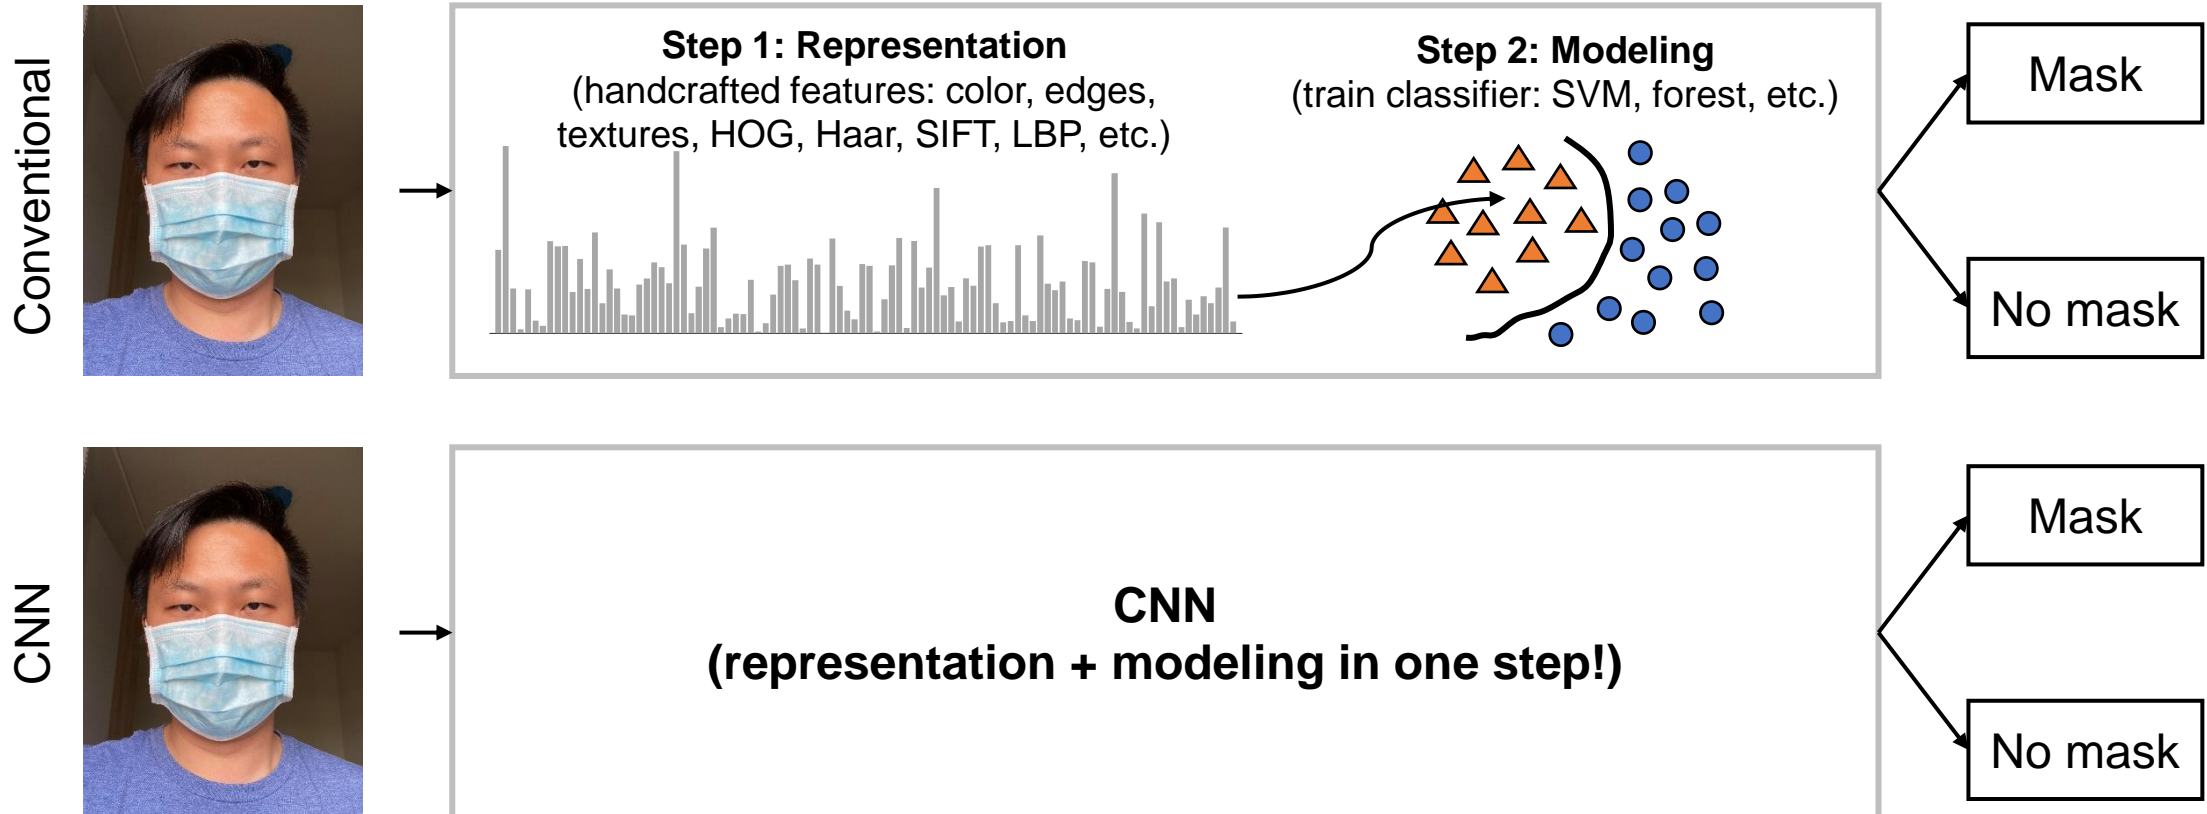


From image data to a decision



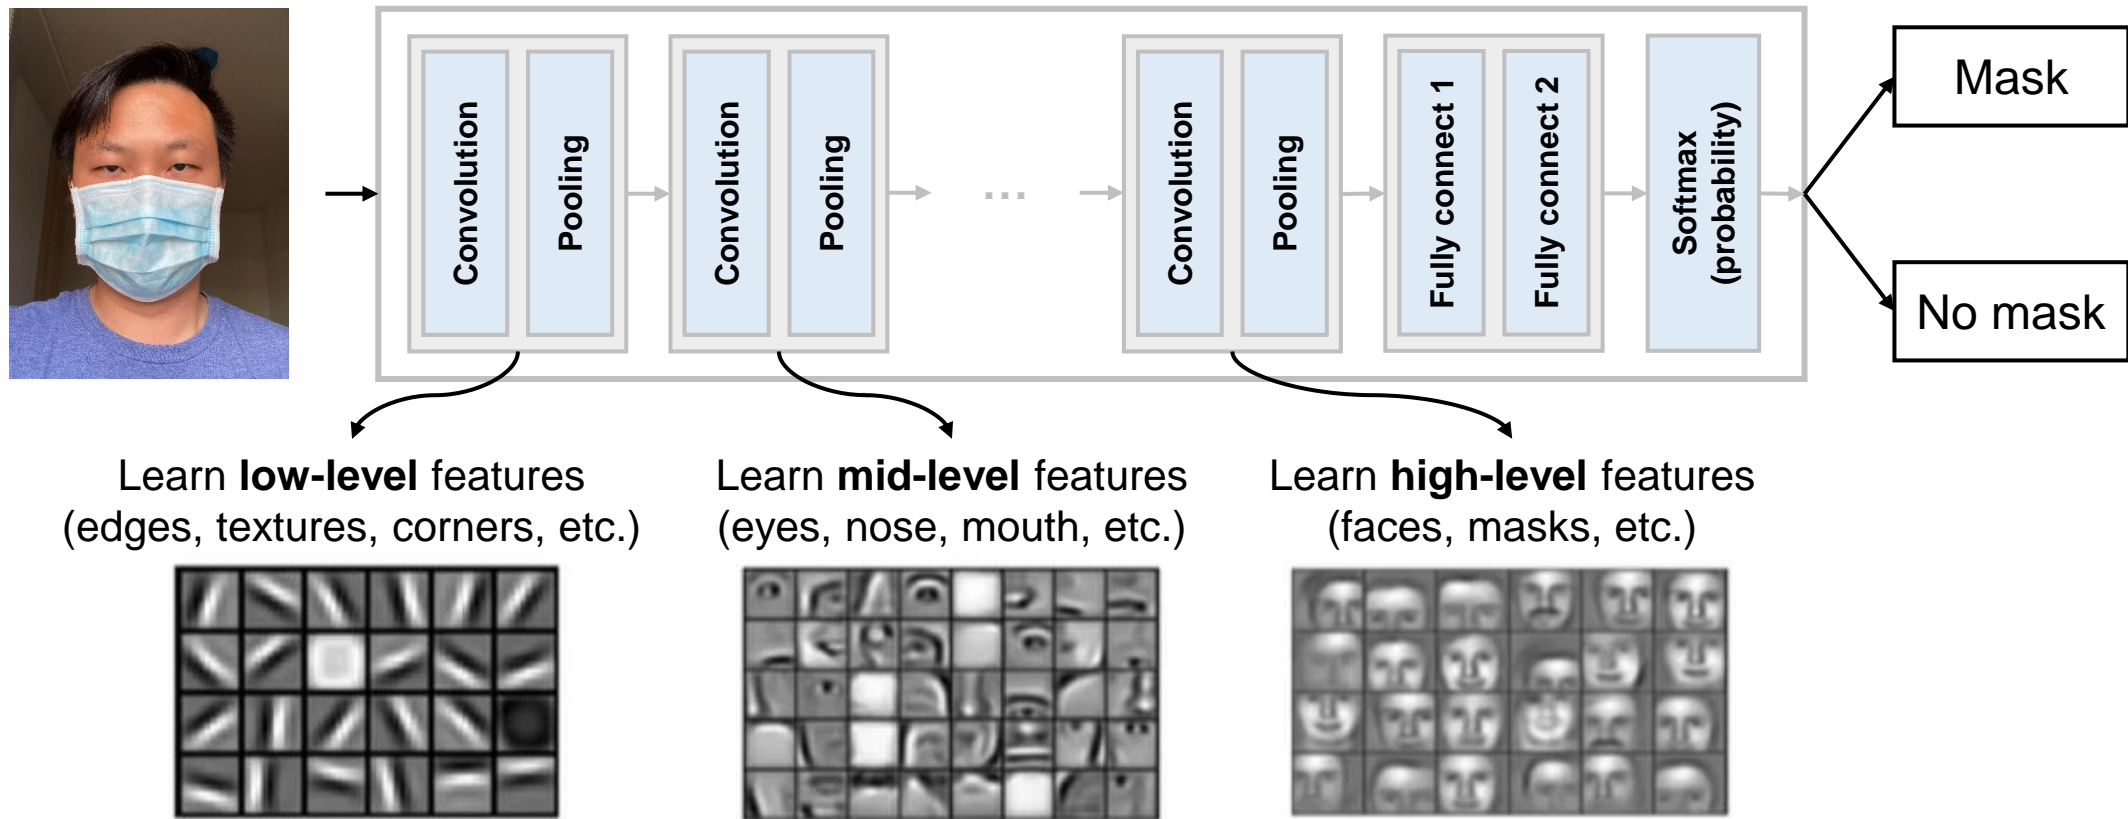
How?

- From conventional machine learning to CNN



What is CNN?

- Overview of a basic CNN architecture



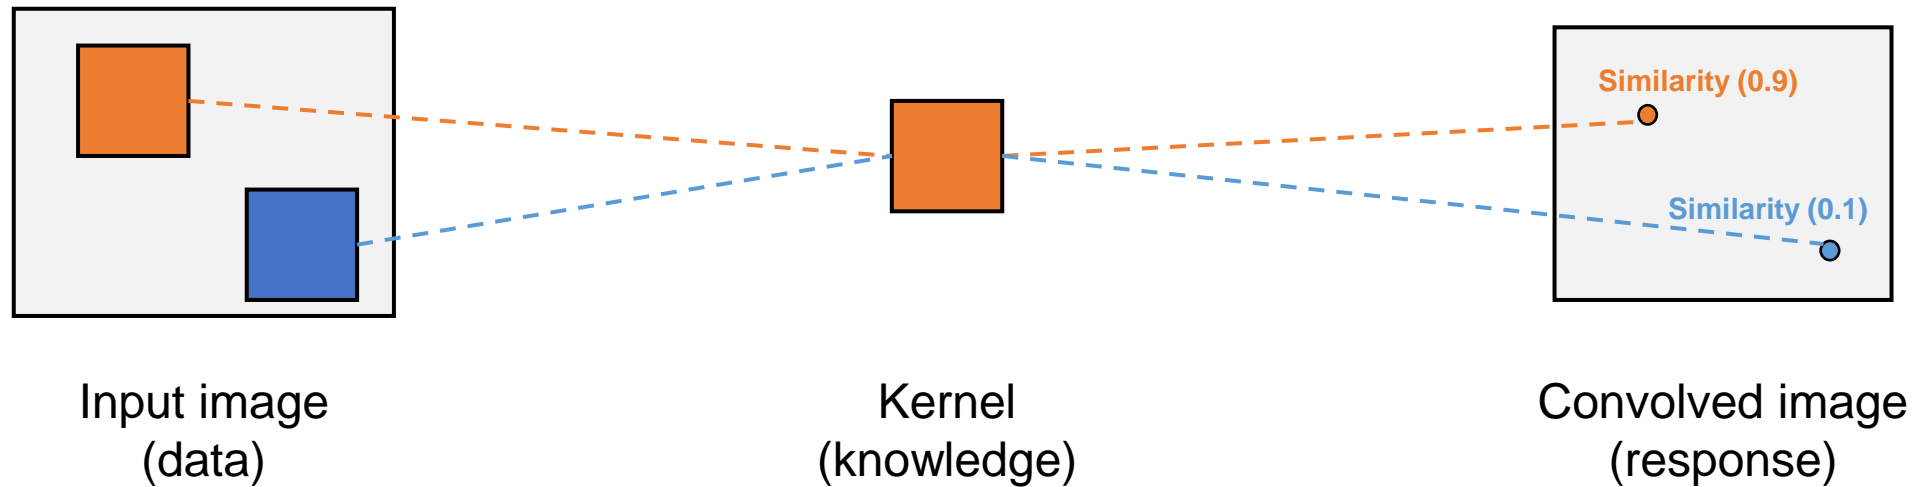
Overview of CNN basics

- Why convolution?
- Convolutional layer
- Pooling layer
- Fully-connected layer
- CNN training
- A minimum CNN example

Why convolution?


Why convolution?

- Convolution in CNN is actually “**cross correlation**” without kernel flipping!



*Essence of cross correlation: **it measures the similarity (inner product) of data and knowledge at different shifted locations.** If data has a pattern similar to the knowledge, it will give a high response (correlated).*

Why convolution?



Why convolution?

- It detects edges!

1	1	1	1	2	2	2	2
1	1	1	1	2	2	2	2
1	1	1	1	2	2	2	2
1	1	1	1	2	2	2	2
3	3	3	3	4	4	4	4
3	3	3	3	4	4	4	4
3	3	3	3	4	4	4	4
3	3	3	3	4	4	4	4

Input image (8x8)

*

-1	-1
1	1

=

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
4	4	4	4	4	4	4	4
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

*

-1	1
-1	1

=

0	0	0	2	0	0	0	0
0	0	0	2	0	0	0	0
0	0	0	2	0	0	0	0
0	0	0	2	0	0	0	0
0	0	0	2	0	0	0	0
0	0	0	2	0	0	0	0
0	0	0	2	0	0	0	0
0	0	0	2	0	0	0	0

Kernels

Convolved images


Why convolution?

- It detects edges!



Input image

*




=



Horiz. edges

*



=




Verti. edges

Kernels

Convolved images


Why convolution?

- It detects objects!




Input image (8x8)

*




=



*



=




Kernels (larger)

Convolved images

Why convolution?

- It detects objects!




Input image

*




=



*



=



Kernels (larger)

Convolved images

Why convolution?


- It detects objects!



*




=



*

Kernel (learned)

=




Input image

Convolved images

Why convolution?


- A single kernel



$6 \times 6 \times 3$

Image size: $n \times n \times c$


*



$3 \times 3 \times 3$

Kernel size: $k \times k \times c$
(stride $(s) = 1$)

=




4×4

$$\left\lceil \frac{n-k}{s} + 1 \right\rceil \times \left\lceil \frac{n-k}{s} + 1 \right\rceil$$

Why convolution?


- Multiple kernels (one is not enough!)



Convolutional layer


Convolutional layer

- Problem: linear operation between layers is redundant!



Convolutional layer


- Solution: add non-linear activation!



Convolutional layer

- What non-linear activations to use?

Sigmoid




$$y = \frac{1}{1 + e^{-x}}$$




Gradient vanishing ($dy \leq 1$)

Tanh



$$y = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

ReLU




$$y = \max(0, x)$$



ReLU is most popular, its derivative is 0 or 1, easy for back propagation training.

Leaky ReLU




$$y = \max(0.1x, x)$$


Convolutional layer

- Why non-linear activation works? Biological reasons?...

Biological neuron



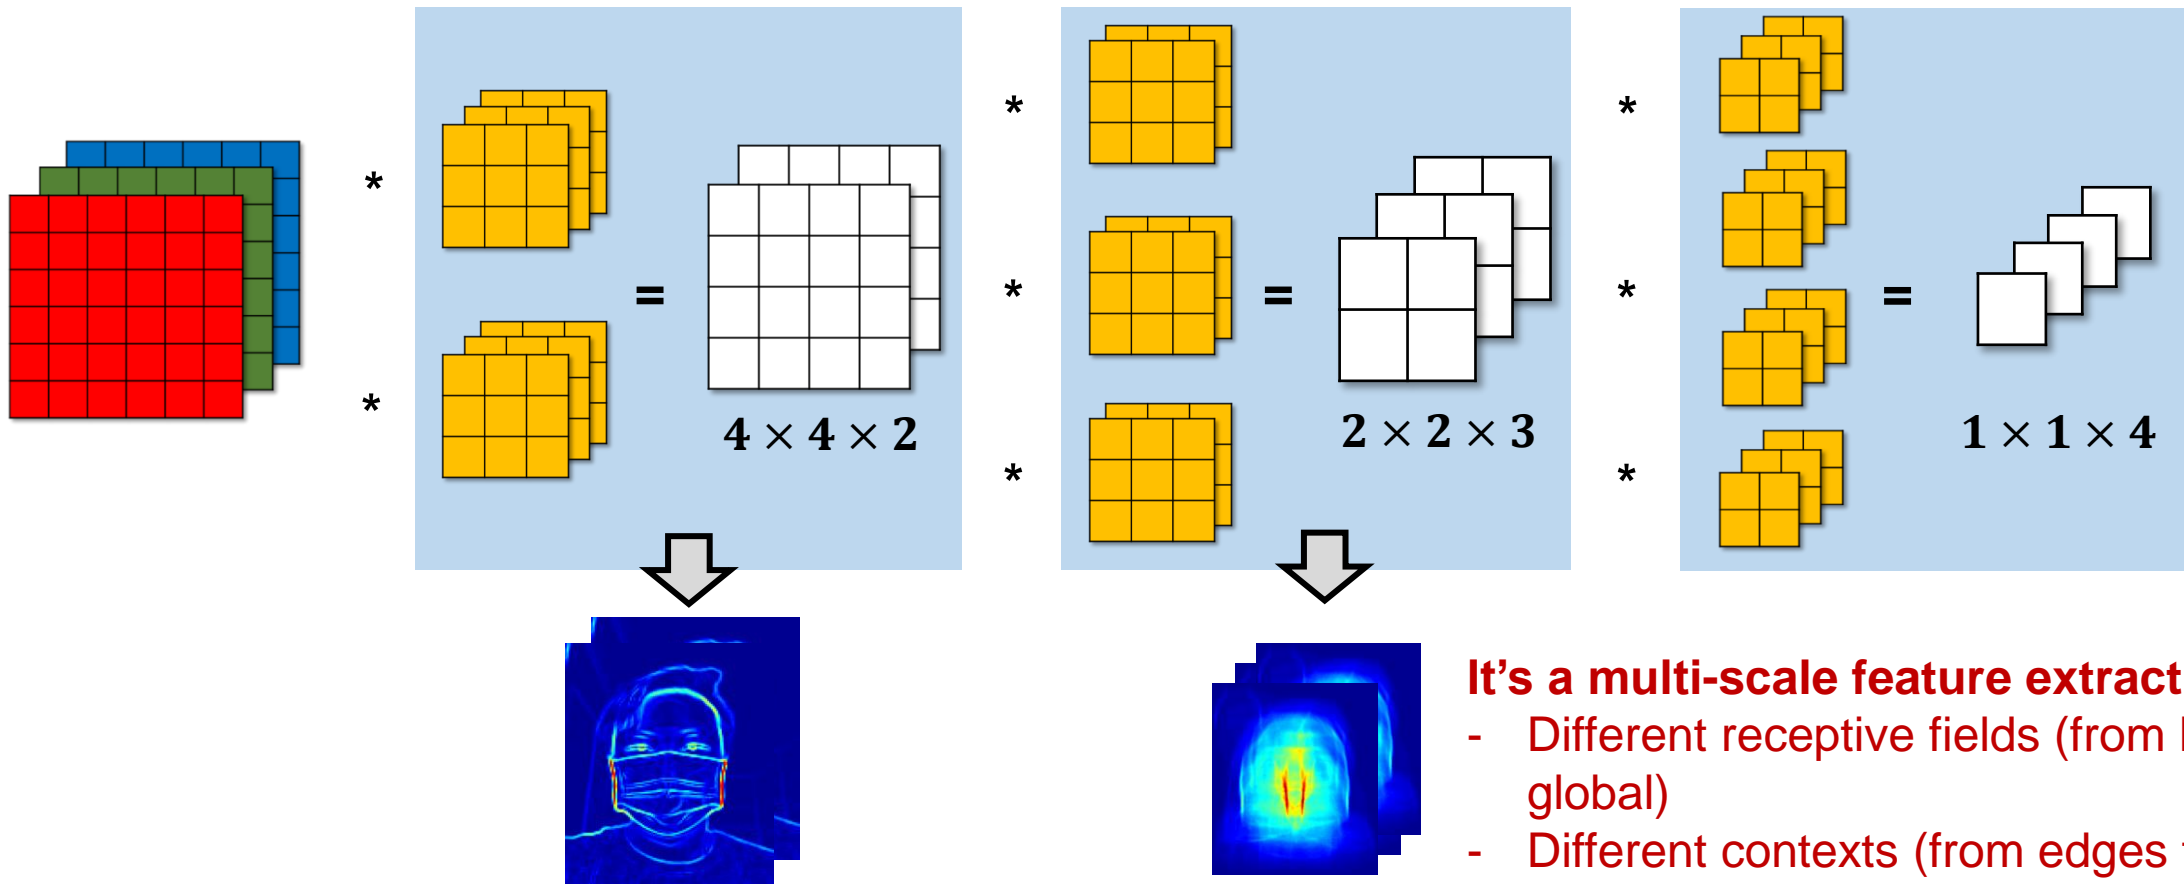
Artificial neuron



“Only when the stimulation is strong enough, it will activate the neuron and transmit the signal. Possibly because our biological system is not linearly reacting to all stimulus. Otherwise our body will be overwhelmed with all kinds of stimulus...”

Convolutional layer

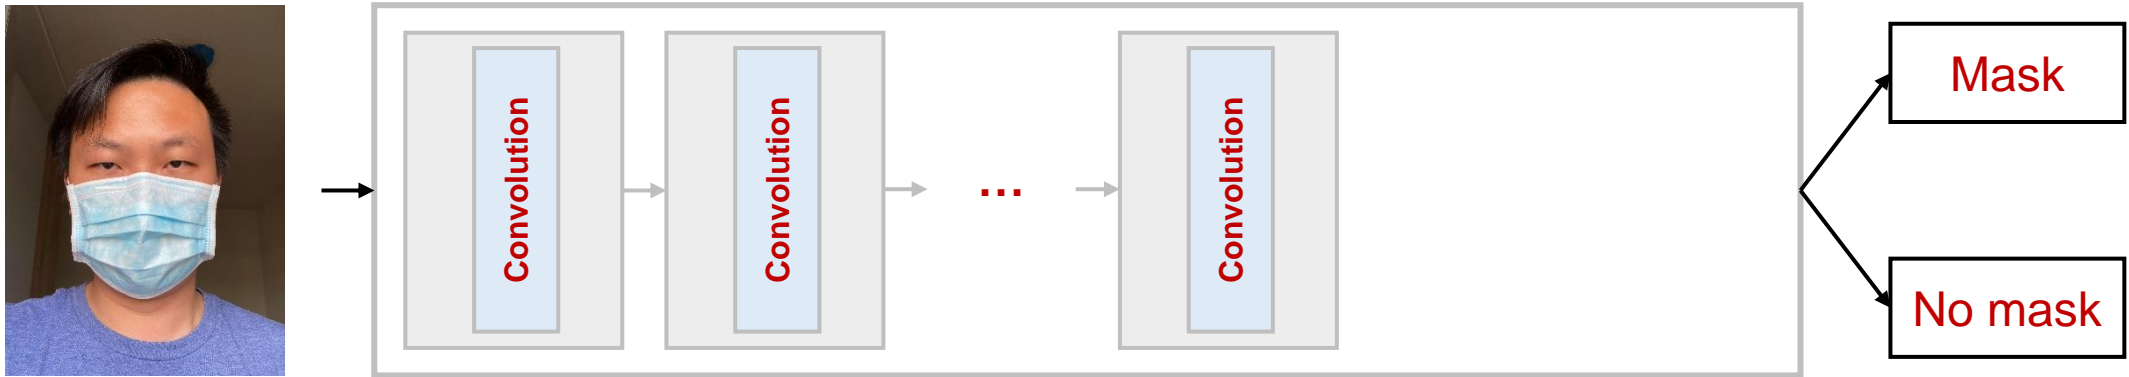
- Multiple convolutional layers: why this is useful?



- It's a multi-scale feature extraction:**
- Different receptive fields (from local to global)
 - Different contexts (from edges to objects)

Convolutional layer (in CNN architecture)


- Now we put **convolutional layers** in CNN



Pooling layer

Pooling layer

- Use spatial pooling to speed up down-sampling!




Max pooling is preferred in most cases:

- Highly non-linear
- Salient/high-frequency information (local maxima)
- Translation invariant

**Pooling does not
require parameters!**

Pooling layer


- Max pooling



More or less
translation
invariant

Pooling layer (in CNN architecture)

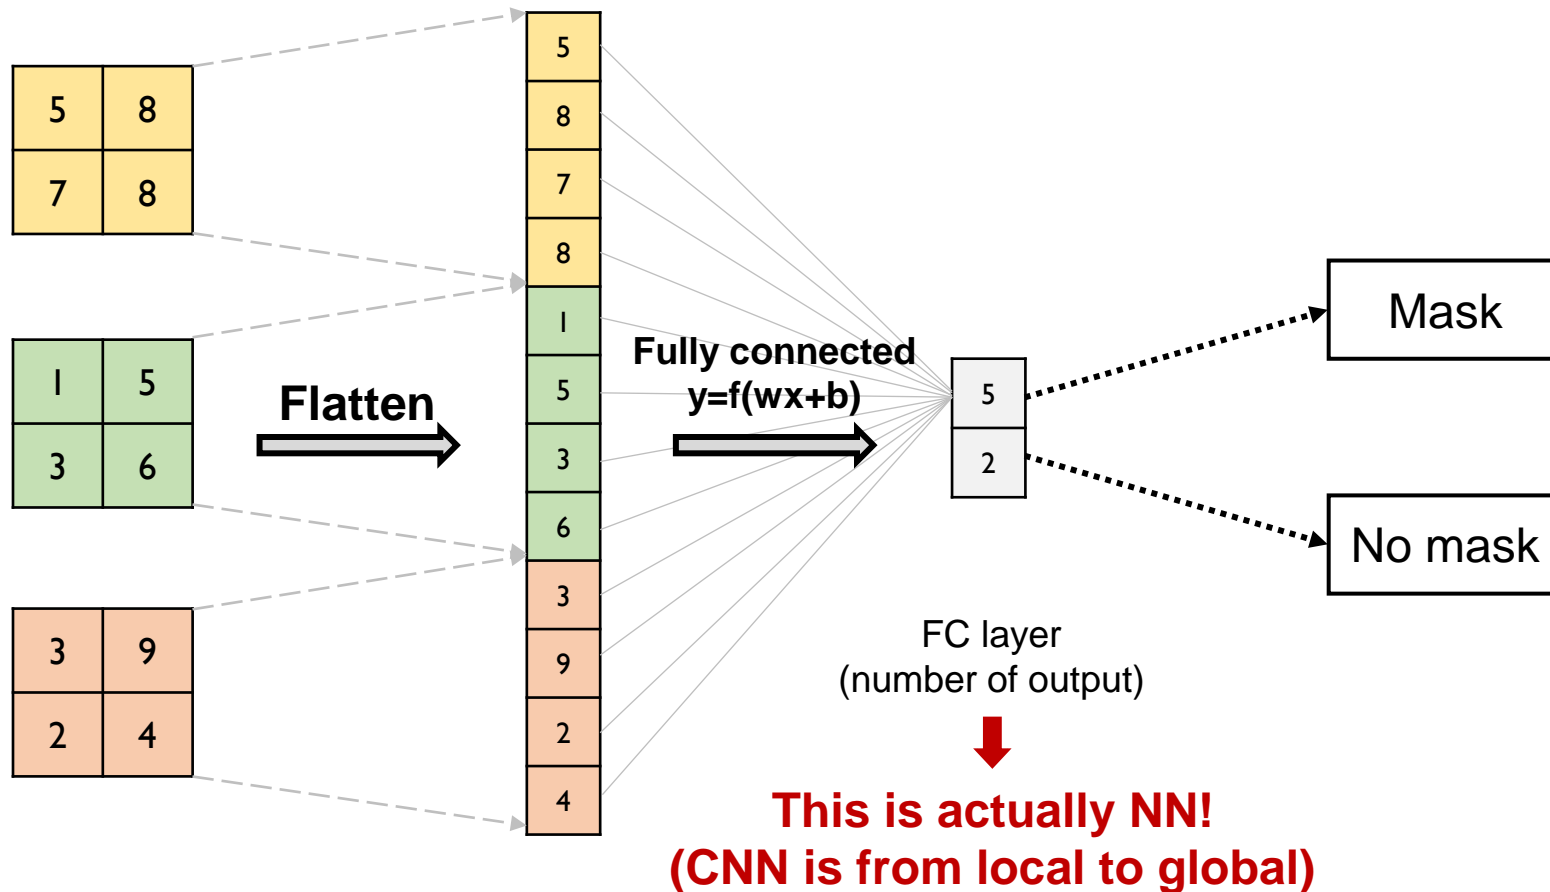
- Now we put **pooling layers** in CNN



Fully-connected layer

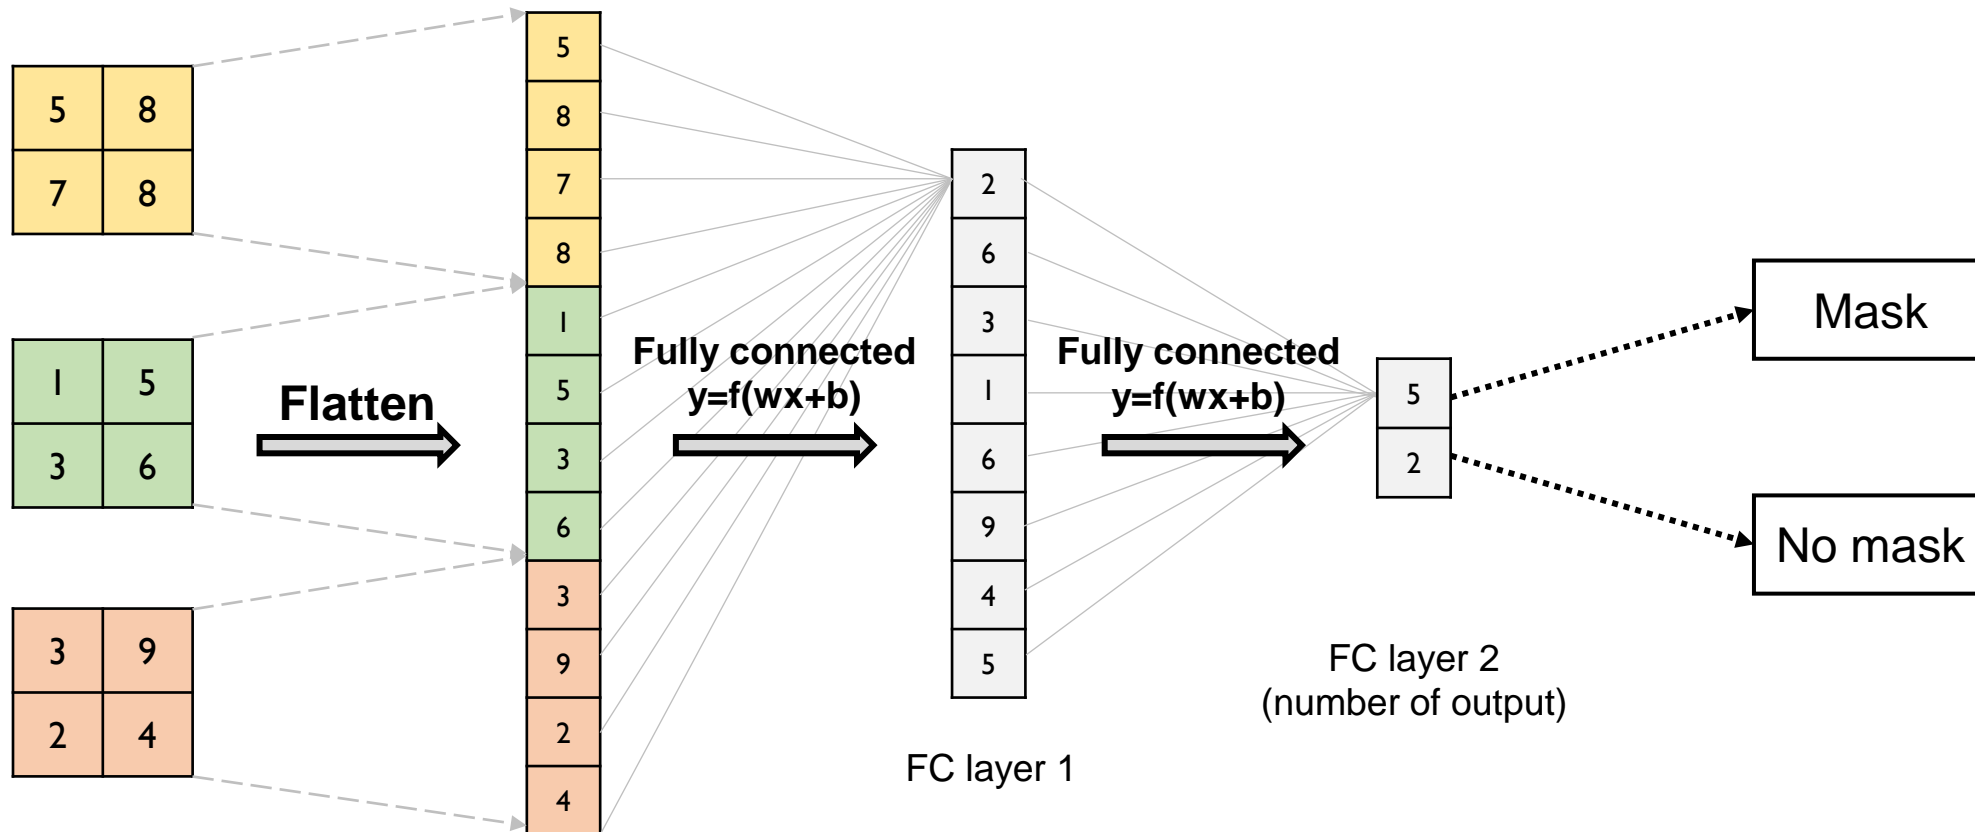
Fully-connected layer

- Convert local 2D maps to a global 1D vector (1 FC layer)



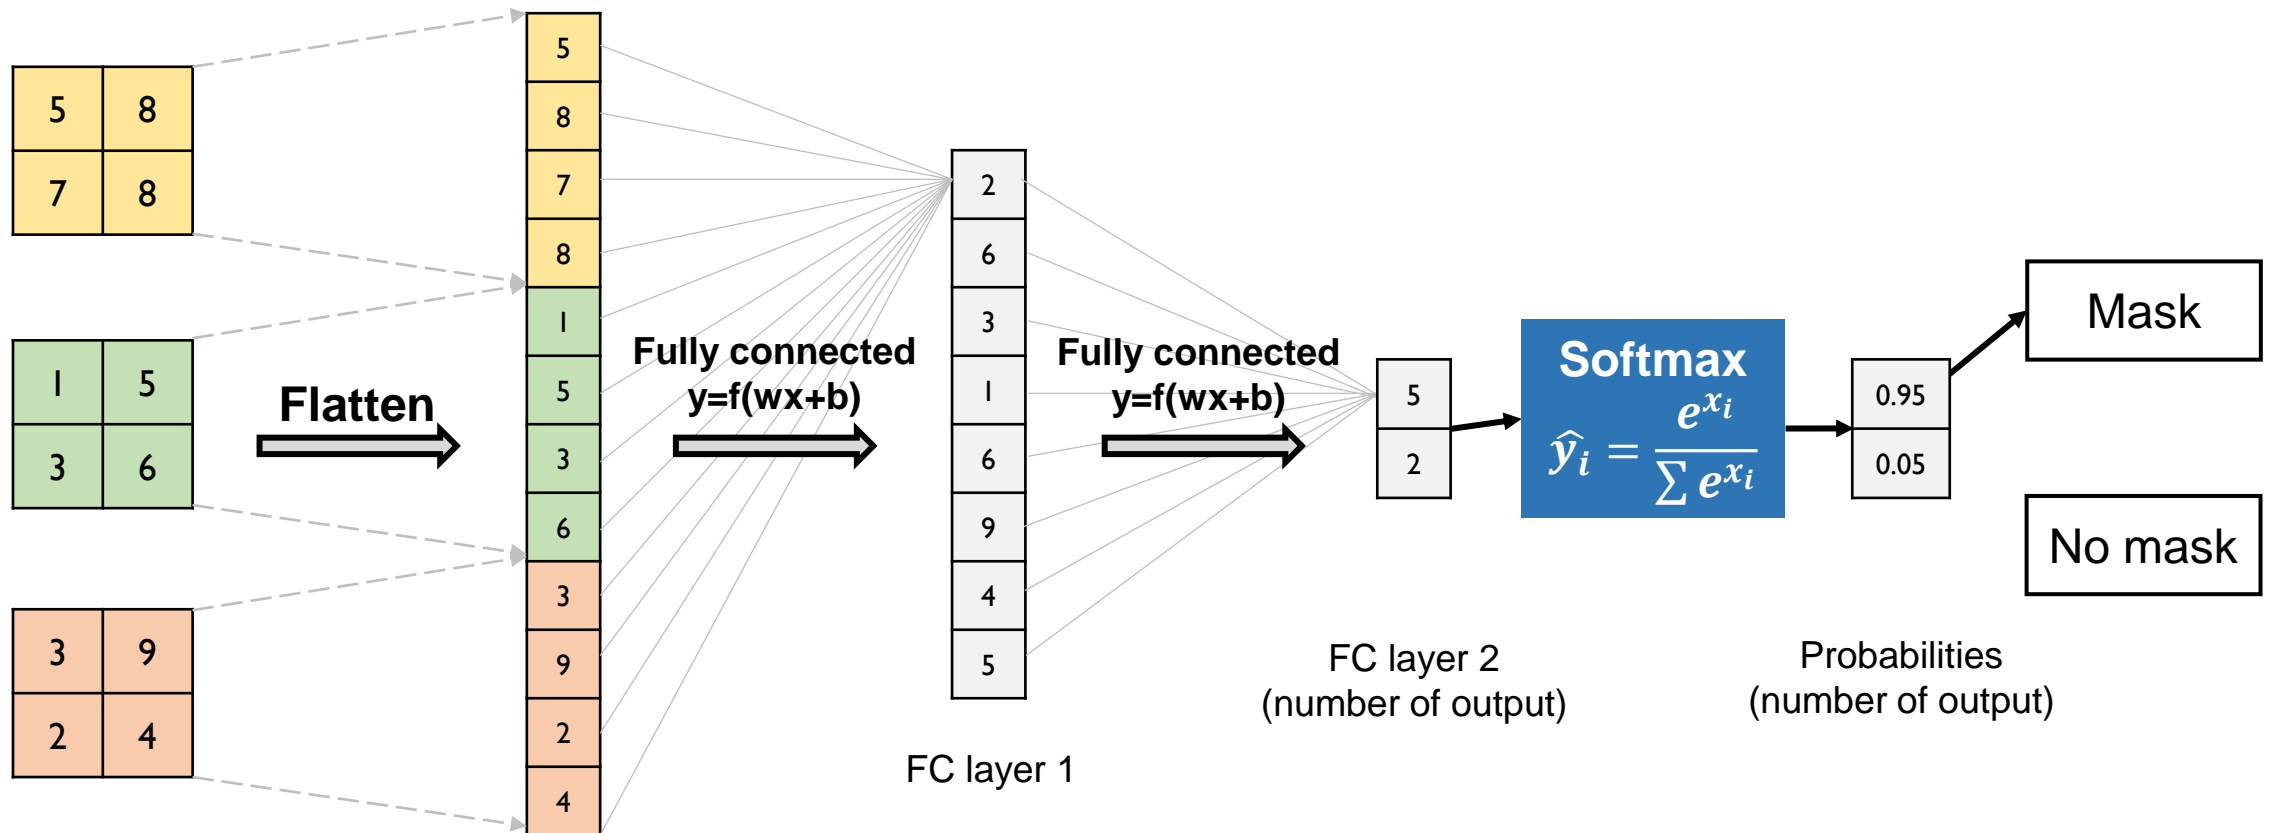
Fully-connected layer

- Convert local 2D maps to a global 1D vector (2 FC layers)



Fully-connected layer


- Convert 1D feature vector to probabilities (Softmax)



Fully-connected layer

- Why fully-connected is useful?

Conventional




Local features

(though features are concatenated, still local)

This is the powerful (terrible?) part of CNN:

- It may use global semantics to make a decision
- It does not even need to see a mask, but can infer a “mask” based on surrounding context
- It gives a global optimal decision, but no idea what is the measurement origin.

CNN




Fully-connected features

(features inside/outside object are fully exploited)

Fully-connected layer (in CNN architecture)

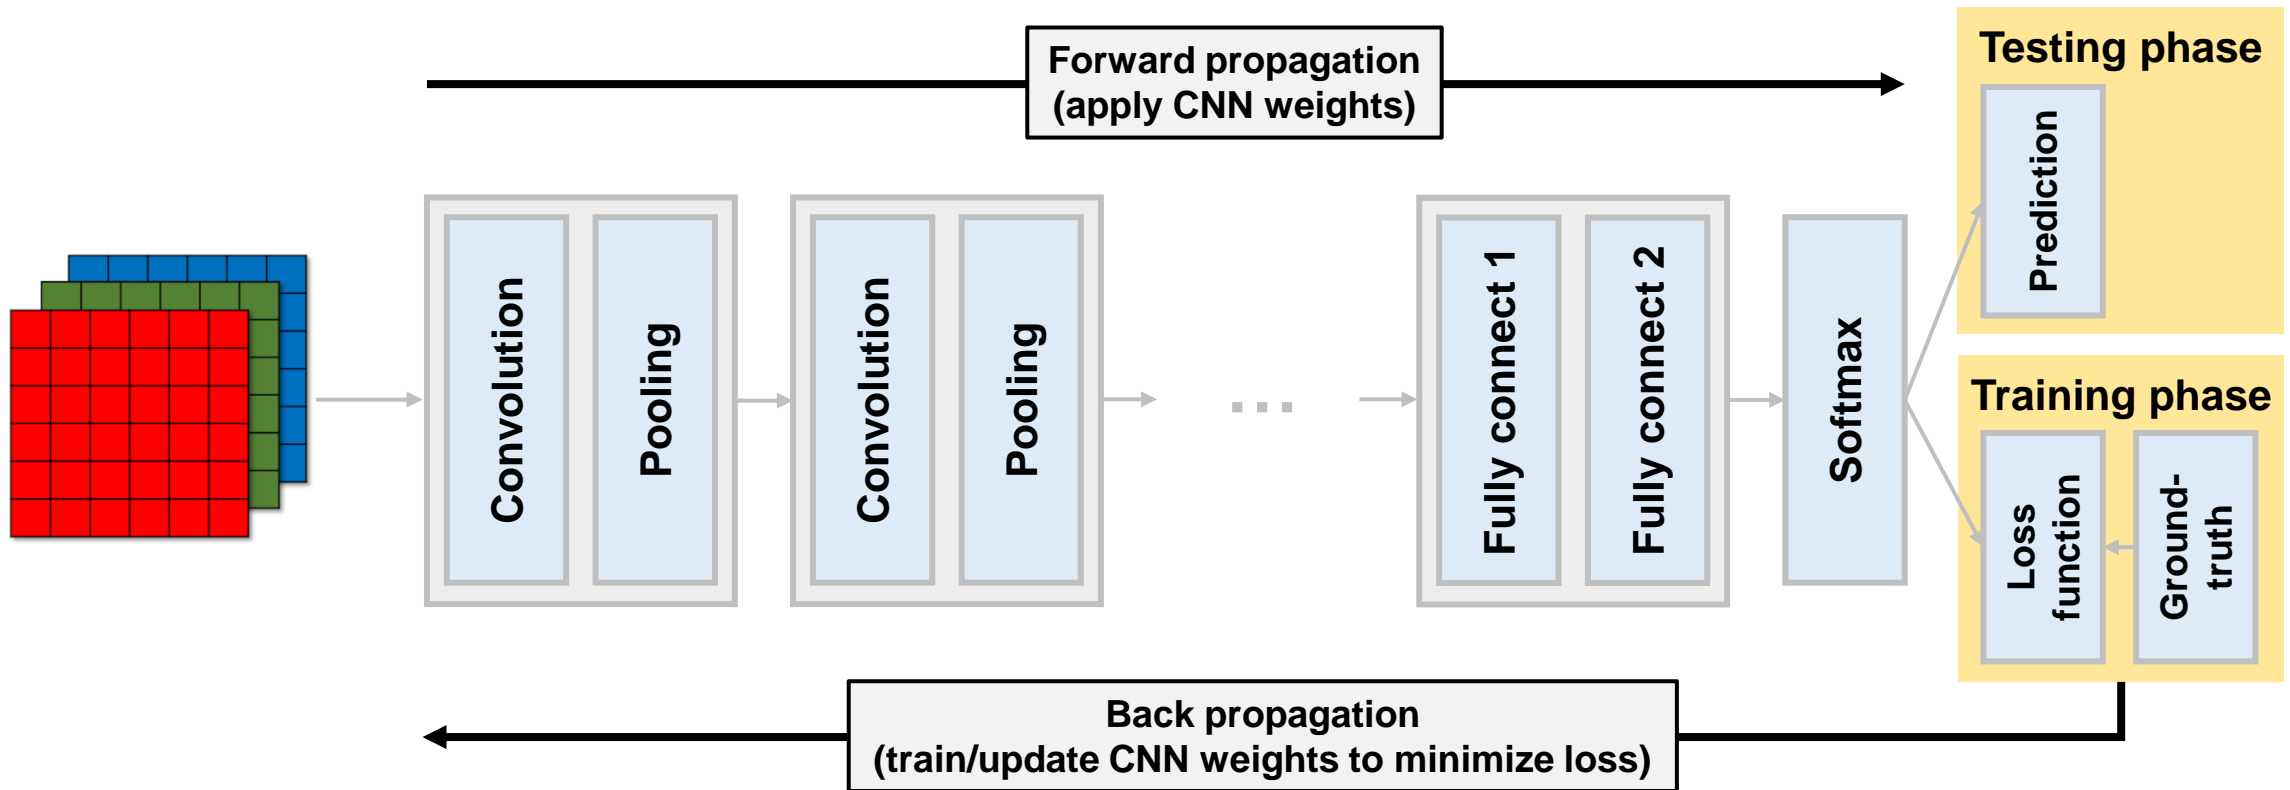
- Now we put **fully-connected layers** in CNN



Application-wise it is more or less ready.
But you do not get the CNN weights for free!

How to get CNN weights?

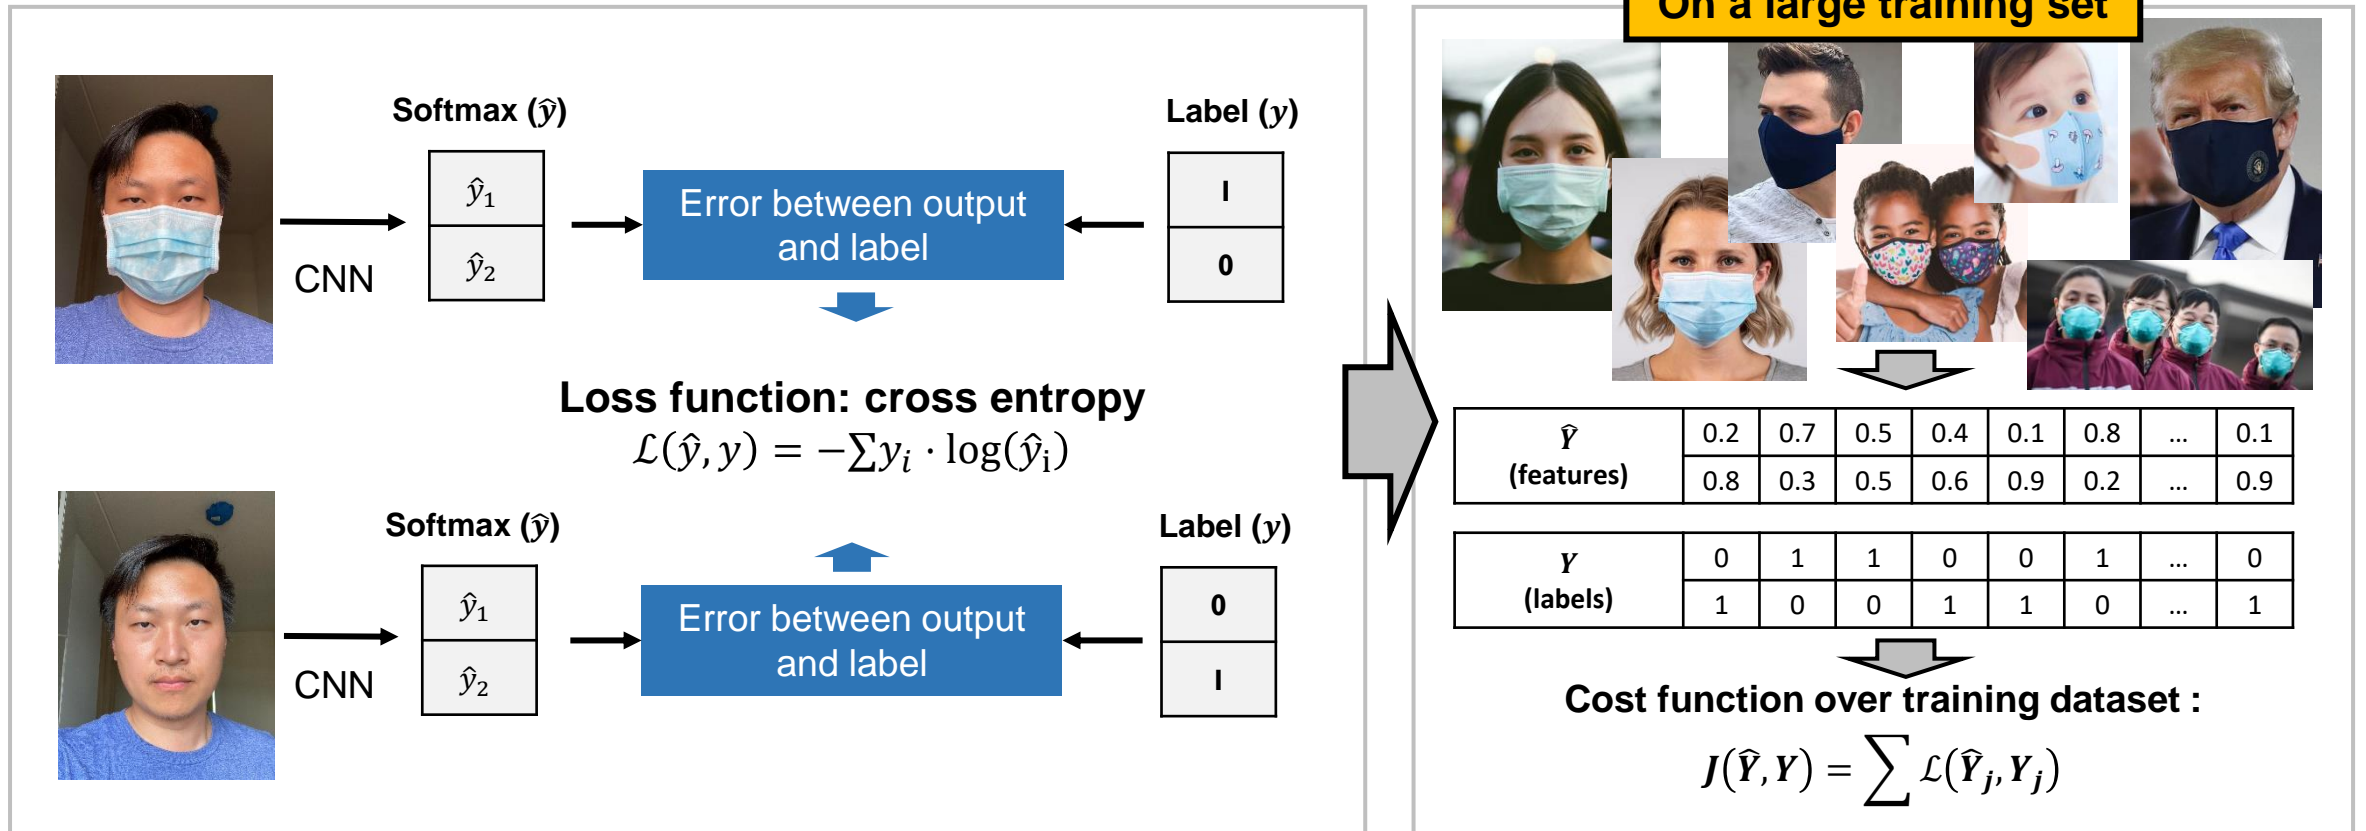
- Train it!!!



CNN training

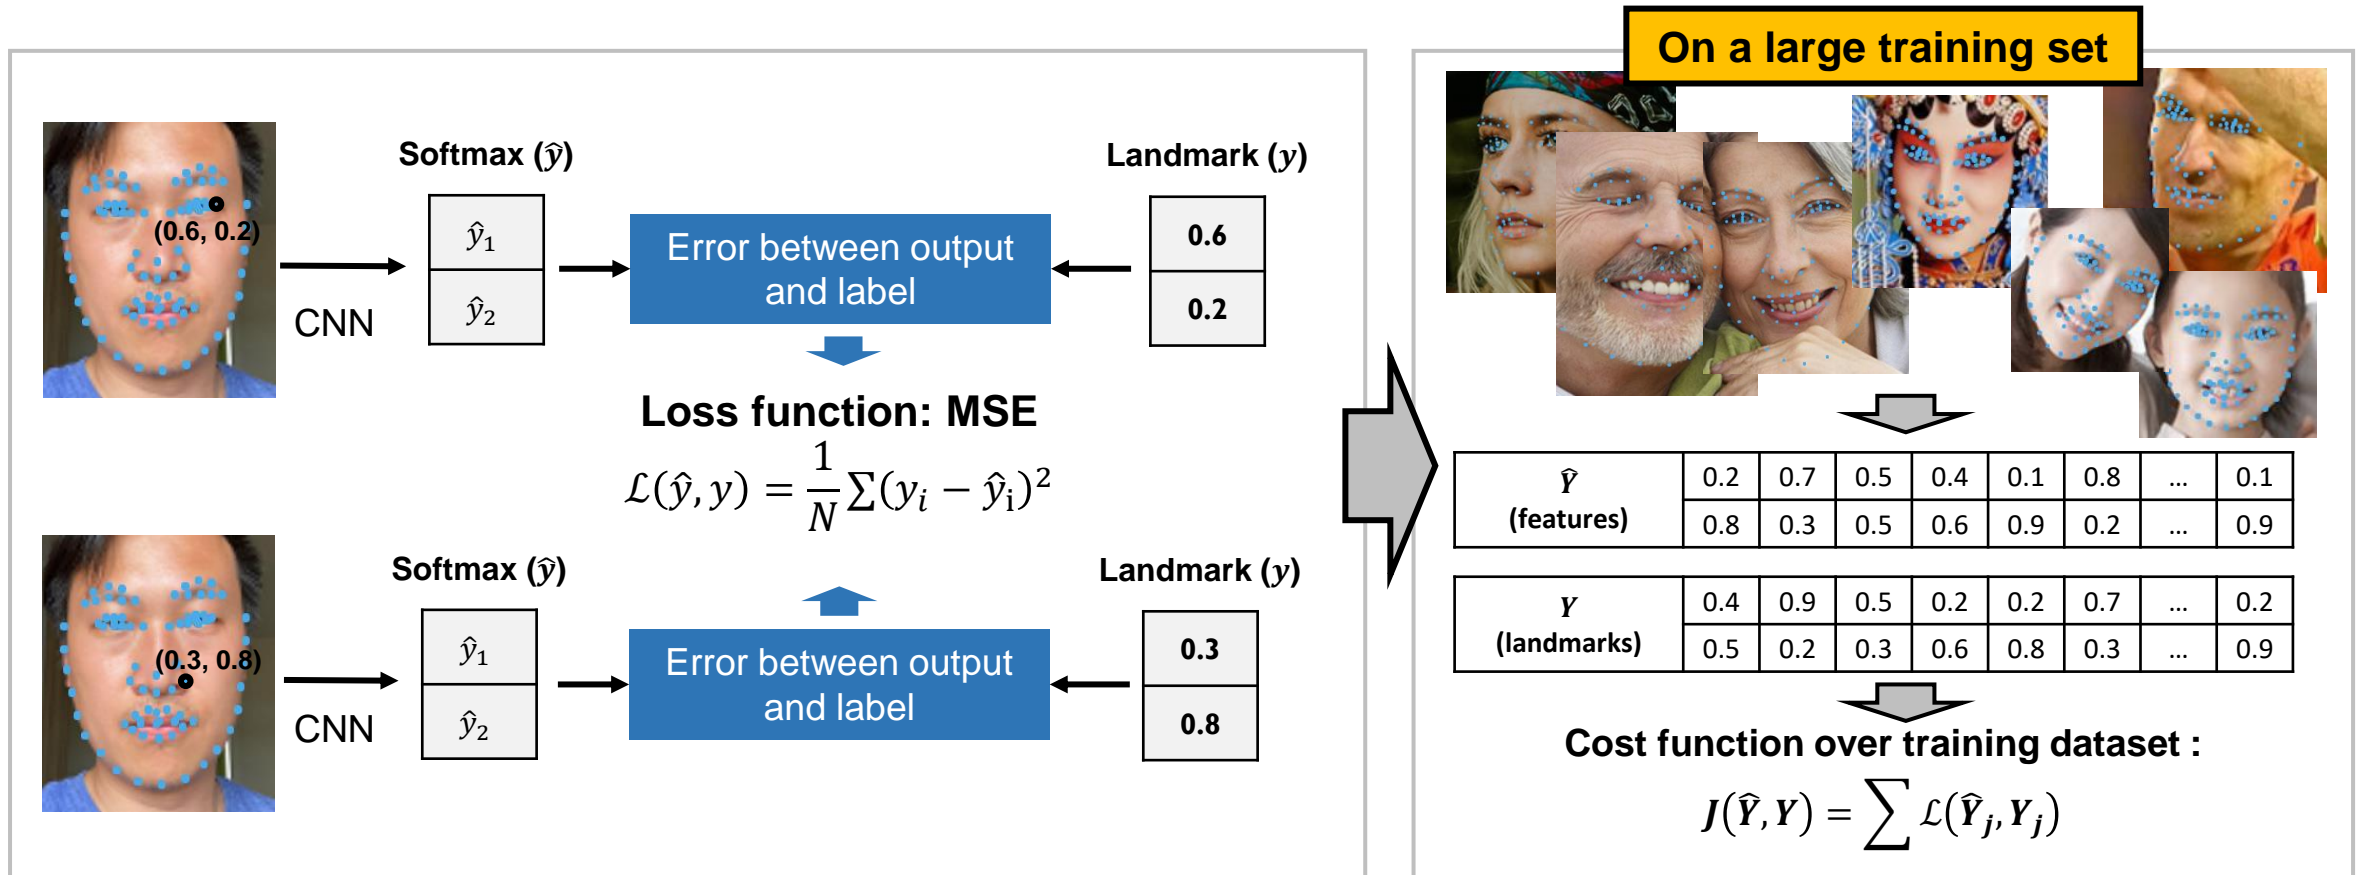
Objective function

- First define an objective function (e.g. classification)



Objective function

- First define an objective function (e.g. regression)



Gradient descent optimization

- Train CNN weights to minimize the loss towards objective!

- **CNN parameters:** $\theta = \{w, b\}$, w=weights, b=bias
- **Loss function:** $\mathcal{L}(f(X, \theta), Y)$, X=data, Y=label

Optimize: $\min_{\theta} \sum \mathcal{L}(f(X, \theta), Y) + \lambda R(\theta)$



Regularization:
constrain
parameters



$\mathcal{L}(\theta)$

Gradient descent

Update θ in gradient direction (fastest):

$$\theta' = \theta - \alpha \frac{d\mathcal{L}(\theta)}{d\theta}$$

Learning rate

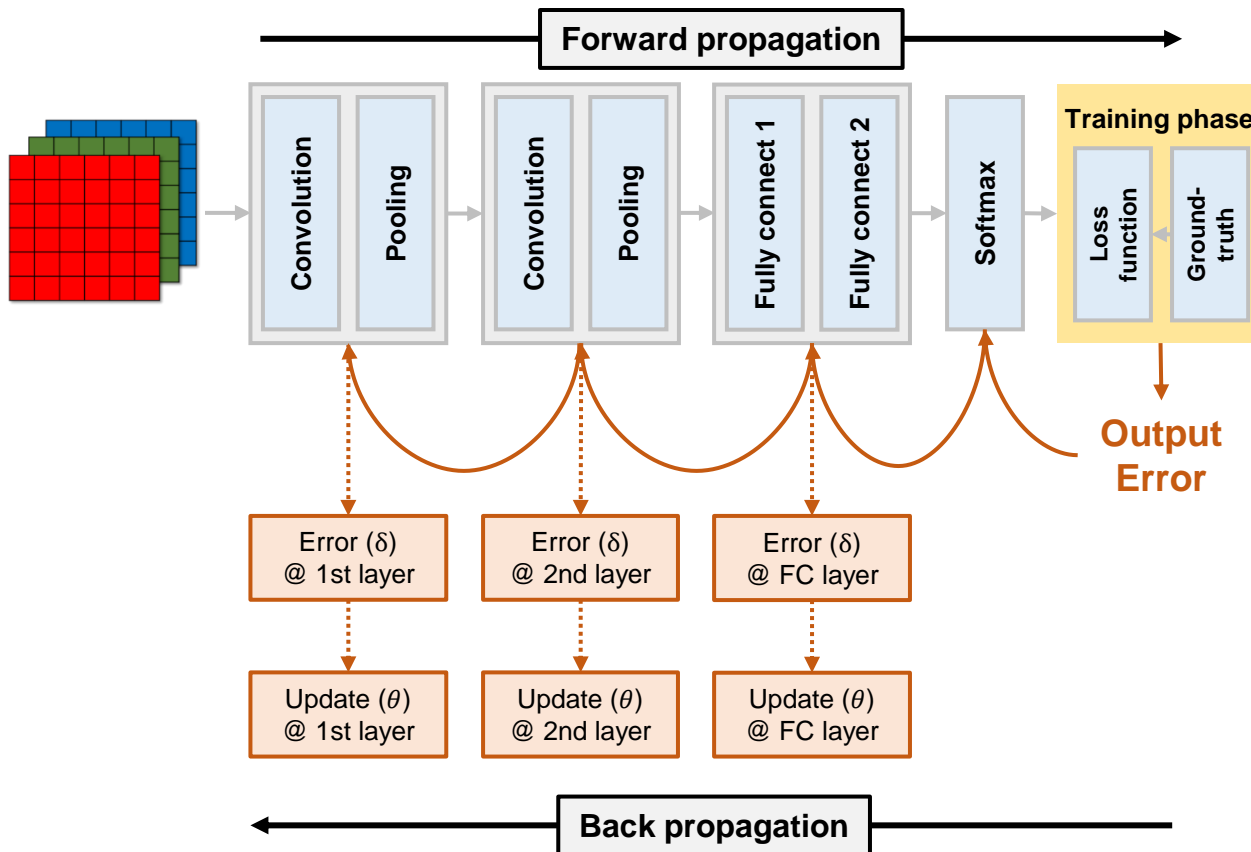
Initial θ

Best θ

θ

Back propagation

- Propagate error back to layers to update weights per layer.



Initialize CNN model (θ)

Do

- Forward propagation
- Back propagation
 - **Estimate output error**
 - **Propagate error back to each layer**
 - **Update θ per layer (gradient descent)**

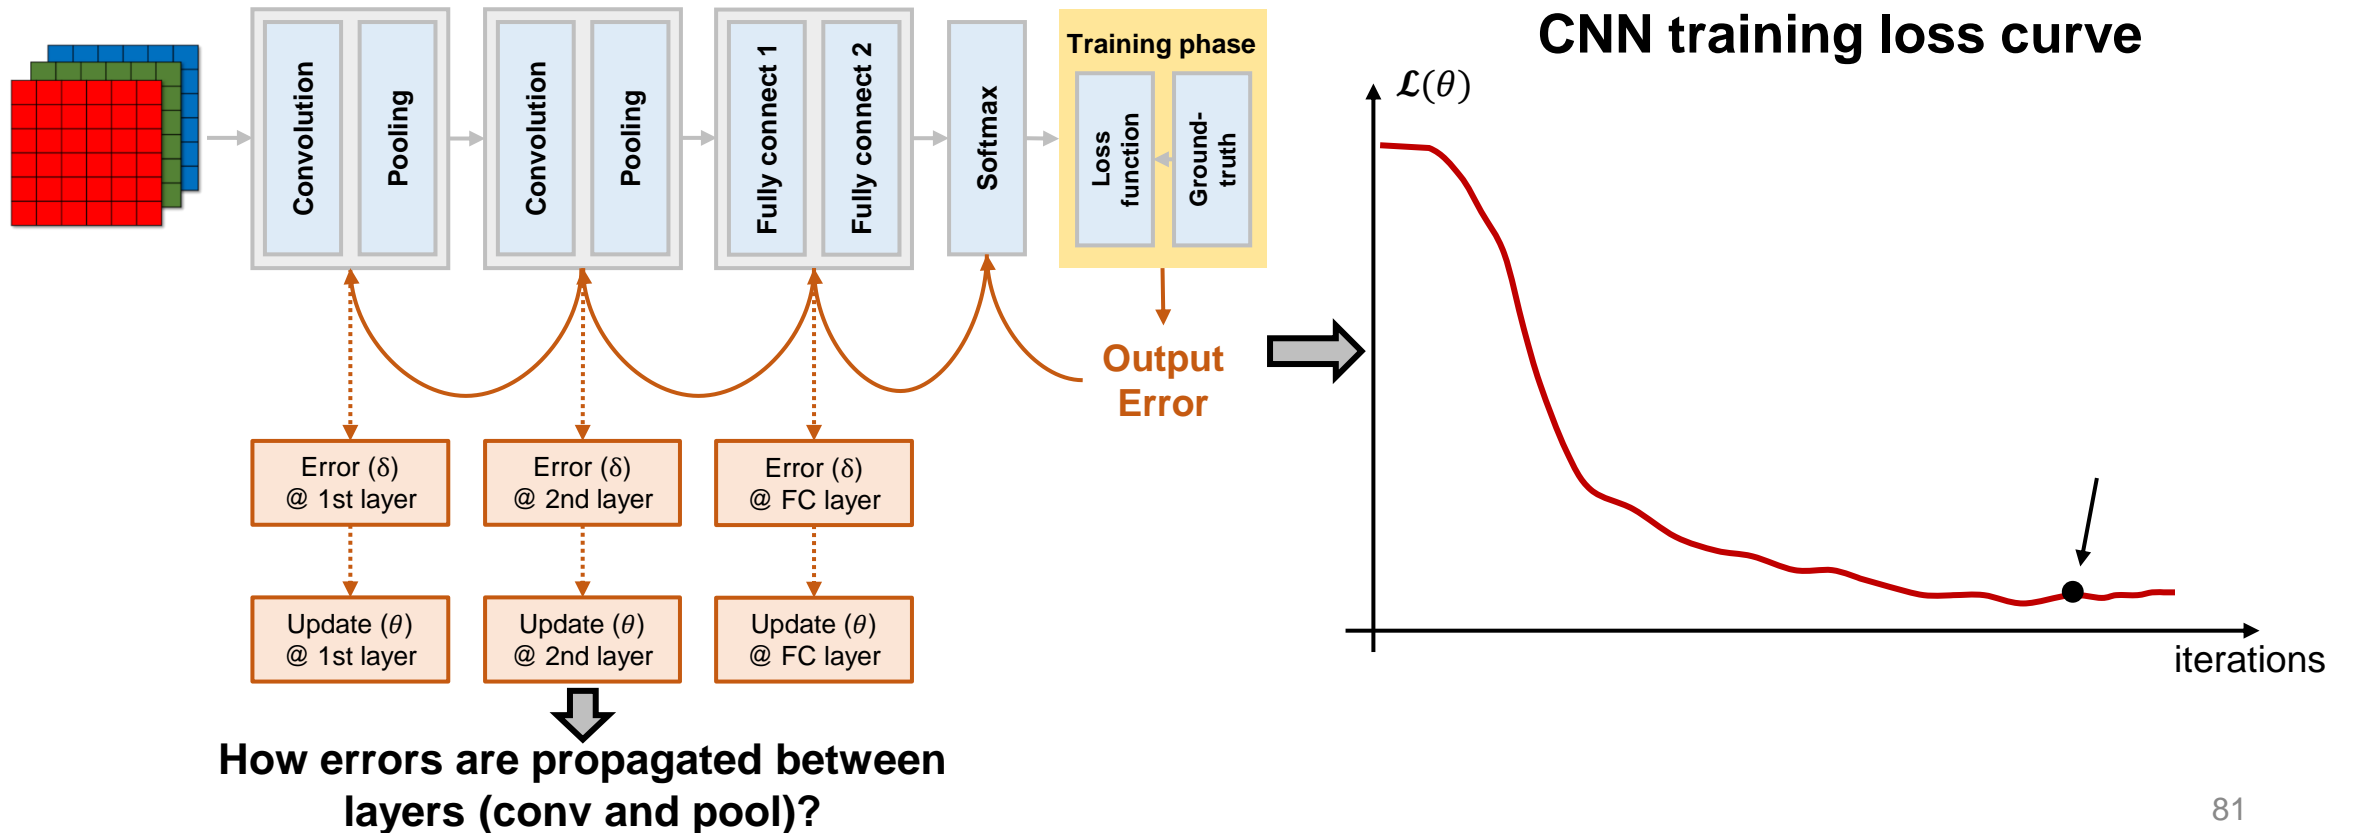
Till loss $\mathcal{L}(\theta)$ is small

Output CNN model (θ)

Chain rule:
$$\frac{dL}{d_{input}} = \frac{dL}{d_{output}} \cdot \frac{d_{output}}{df(wx+b)} \cdot \frac{df(wx+b)}{d_{input}}$$

Back propagation

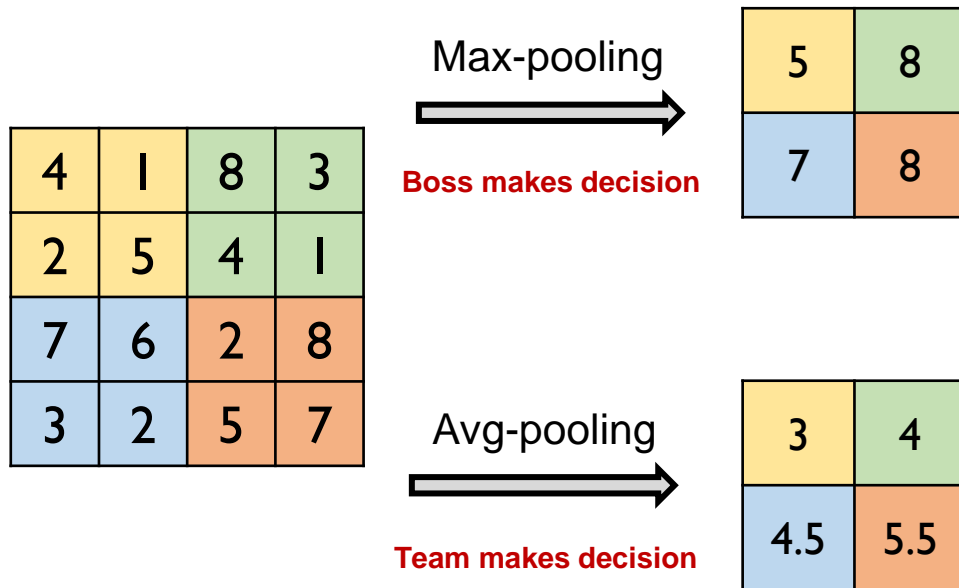
- Propagate error back to layers to update weights per layer.



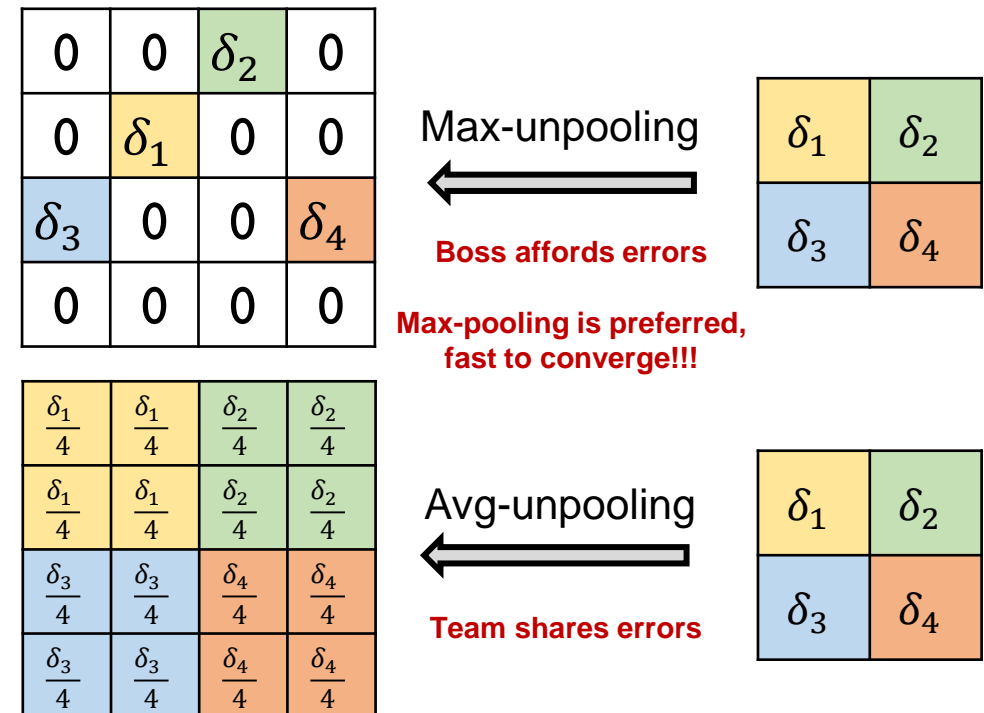
Back propagation

- Back propagation for pooling layer

Pooling (down-sampling)





Unpooling (up-sampling)



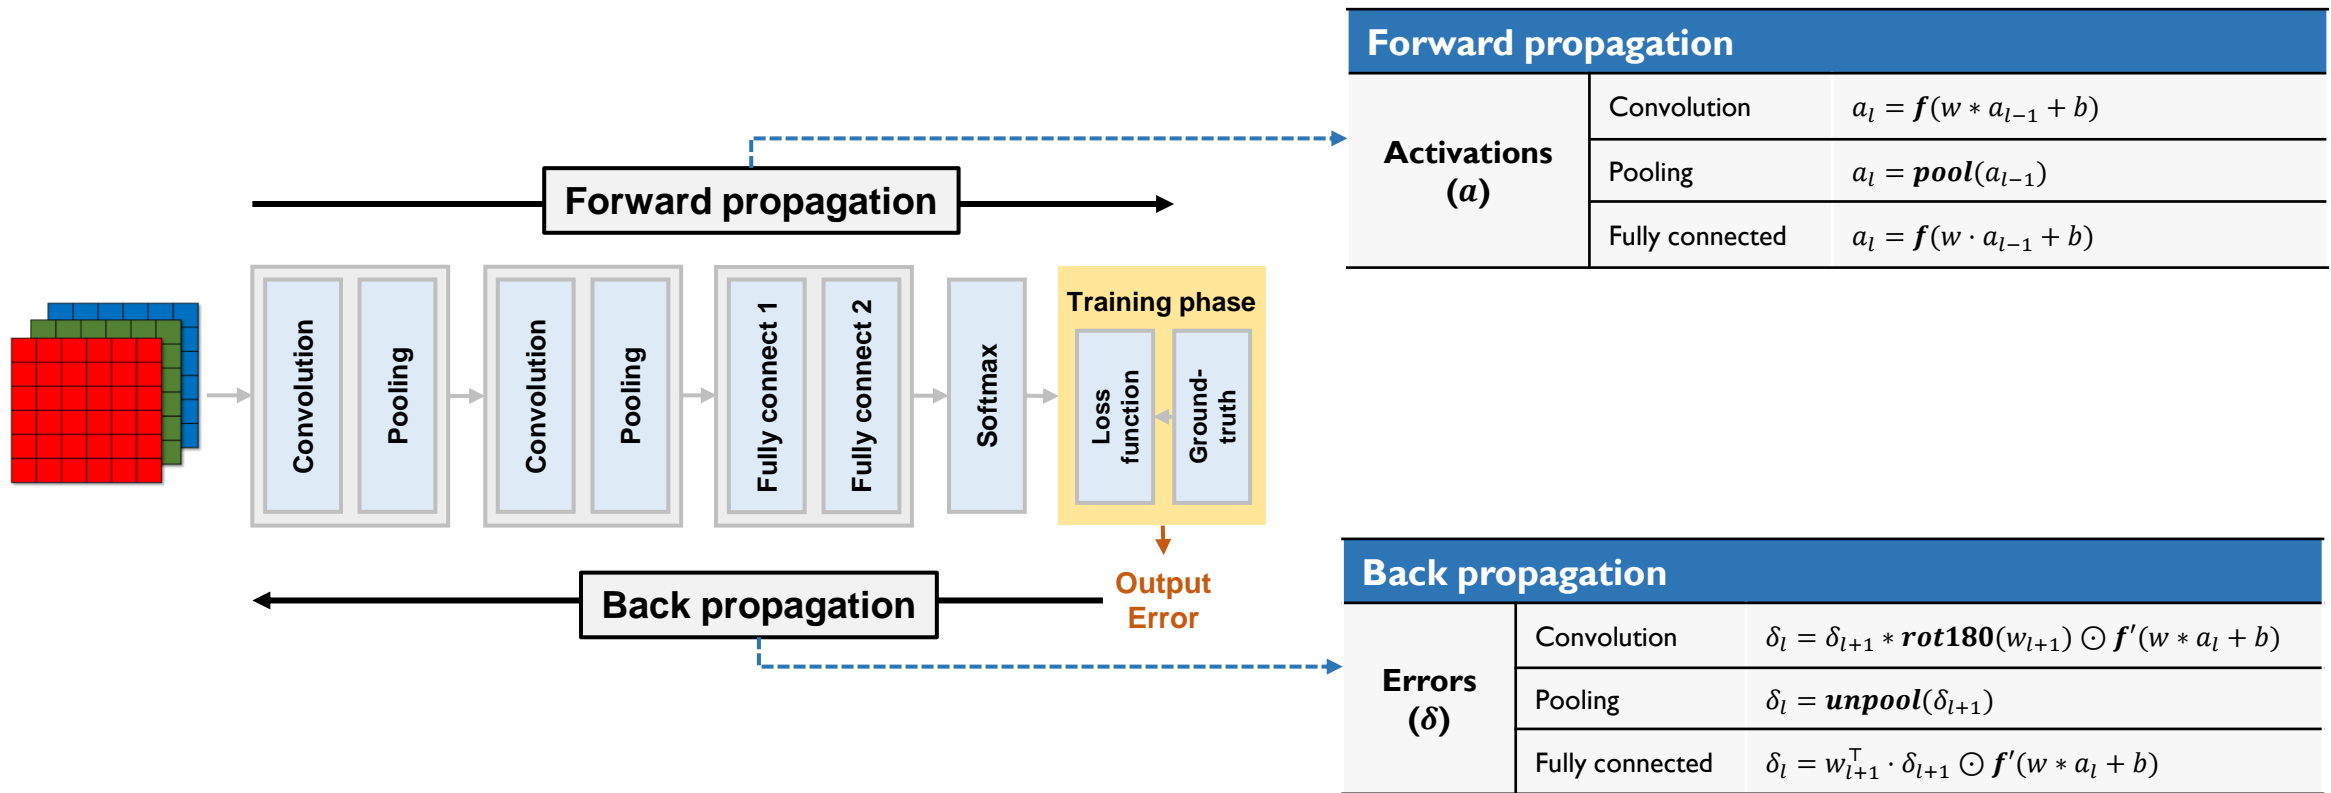
Back propagation

- Back propagation for convolution layer



Back propagation

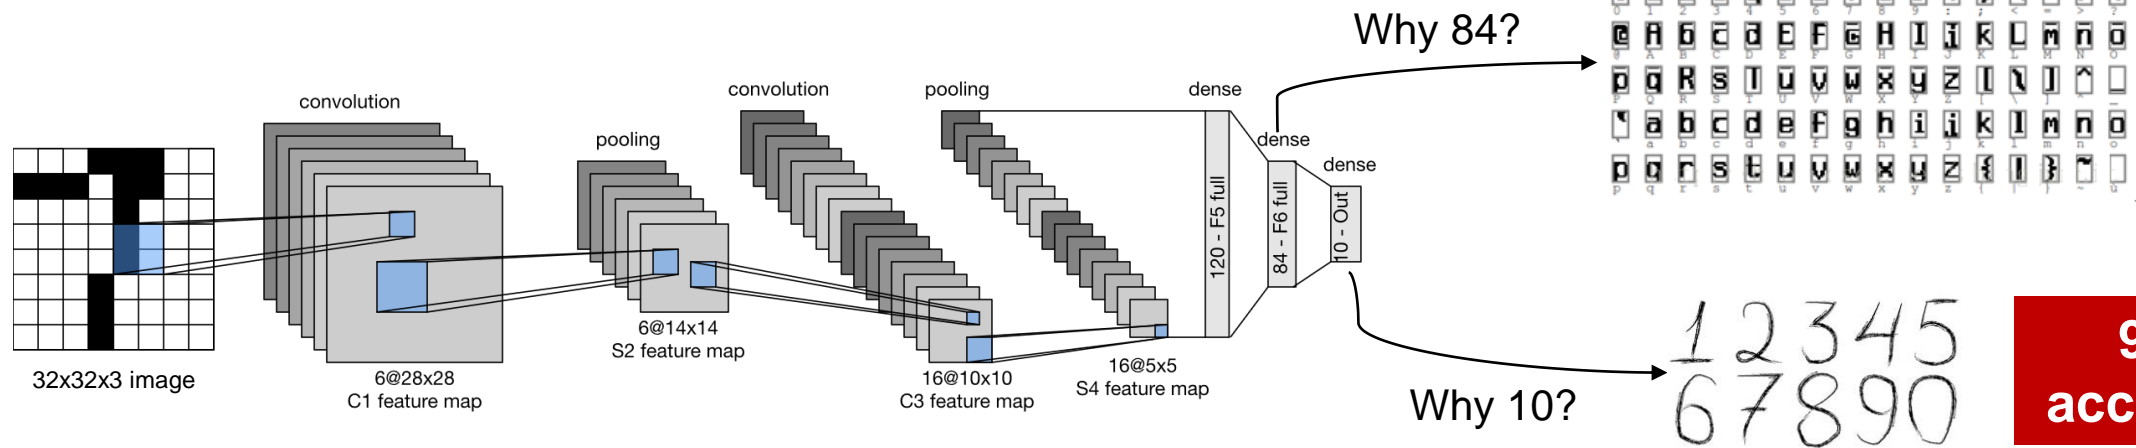
- 6 core equations you need to know for CNN!



A minimal CNN example

LeNet-5

- LeCun et al. (1998) for recognizing handwritten digits

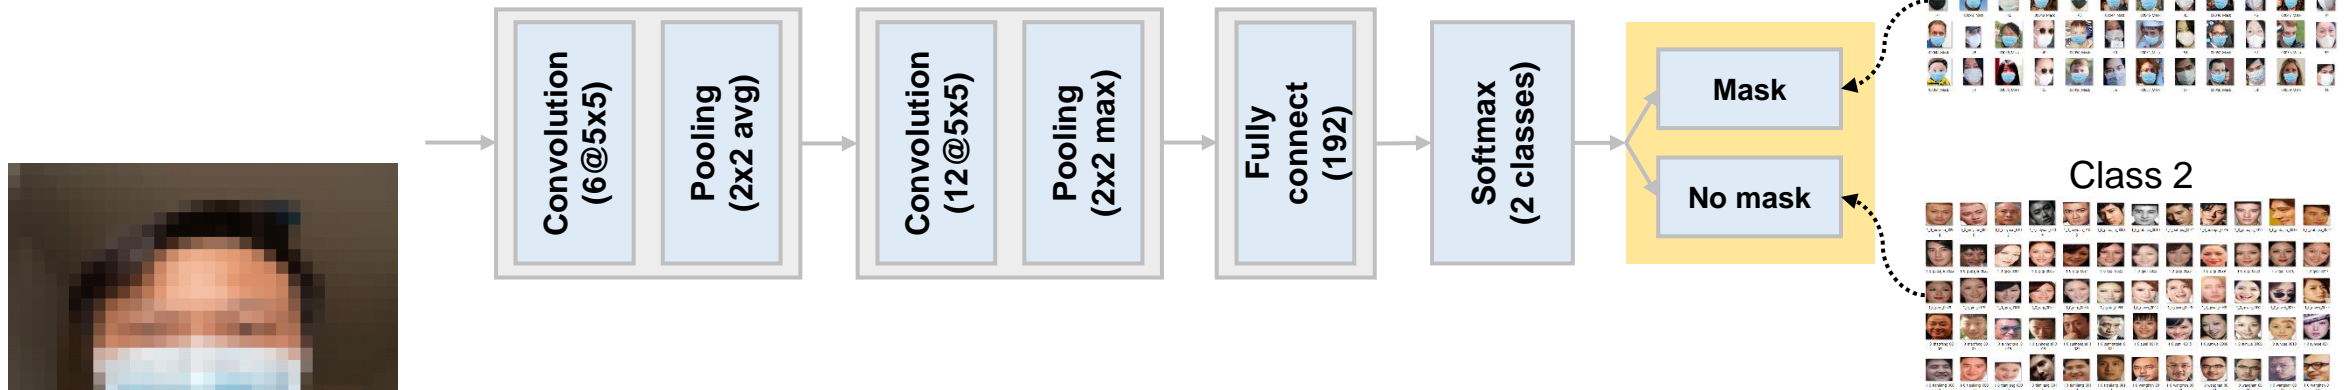


Layers	1	2		3		3			4
Type	Normalization	Convolution	Pooling (2x2)	Convolution	Pooling (2x2)	FC 1	FC 2	FC 2	Softmax
Input size	32x32	32x32	28x28x6	14x14x6	10x10x16	5x5x16	1x1x120	1x1x84	1x1x10
Kernel (weights and bias)	-	6@(5x5+1) = 456	-	16@(5x5x6+1) = 2416	-	120@(5x5x16+1) = 18120	84@(1x1x120+1) = 10164	10@(1x1x84) = 840	-
Output size	32x32	28x28x6	14x14x6	10x10x16	5x5x16	1x1x120	1x1x84	1x1x10	1x1x10


My example

- Implementation from scratch

Demo 1: 2-layers CNN







Demo 2: 1-layer CNN (slightly better)



My example (live demo)

- Very simple, everyone can do





Why CNN works?

Our world is continuous and smooth at different locations and scales.

CNN pitfalls

Attacks for CNN



Cup(16.48%)
Soup Bowl(16.74%)



Bassinet(16.59%)
Paper Towel(16.21%)



Teapot(24.99%)
Joystick(37.39%)




Hamster(35.79%)
Nipple(42.36%)

Just add one pixel or “invisible” noise perturbation can fool deep CNN!



“panda”
57.7% confidence

+ ϵ




=



“gibbon”
99.3% confidence

Attacks for CNN



Adversarial t-shirt, protect privacy?!

Filters (1)

Category

- All Categories
- Clothing
 - Dresses
 - Leggings
 - Mini Skirts
 - Sweatshirts & Hoodies
 - T-Shirts
 - Tank Tops

Gender

- ☒ Everyone
- ☐ Women
- ☐ Men

Style

- ☒ All Styles
- ☐ Active

Adversarial T-Shirts 20 Results

Most Relevant

Product Name	Price
Adversarial Anti-Facial Recogniti... By el-em-cee	€18.29
Adversarial Patch Classic T-Shirt By Dave Sag	€18.81
GAN waking up Tri-blend T-Shirt By clearwhale	€27.09
Adversarial Anti-Facial Recogniti... By el-em-cee	€18.29
Adversarial Sleeveless Top By REApparelCo	€29.54
Adversarial Essential T-Shirt By keygen	€21.11
Adversarial Anti-Facial Recogniti... By el-em-cee	€18.29
Adversarial Anti-Facial Recogniti... By el-em-cee	€18.29
Adversarial Anti-Facial Recogniti... By el-em-cee	€18.29
Adversarial Anti-Facial Recogniti... By el-em-cee	€18.29

CNN remains a black box

- Though we have certain ways/tools to visualize and explain CNN retrospectively, it is not fully transparent and explainable, especially for training process
- No good guidelines to train CNN for a specific task (based on empirical settings or feelings, rule of thumb)
- Hurdle for some assignments, i.e. FDA clearance for clinical device
- ...

CNN changes the way of research


Classical research style
(compete on ideas and domain knowledge)



Research becomes boring?!



CNN research style
(compete on data collection and GPU)



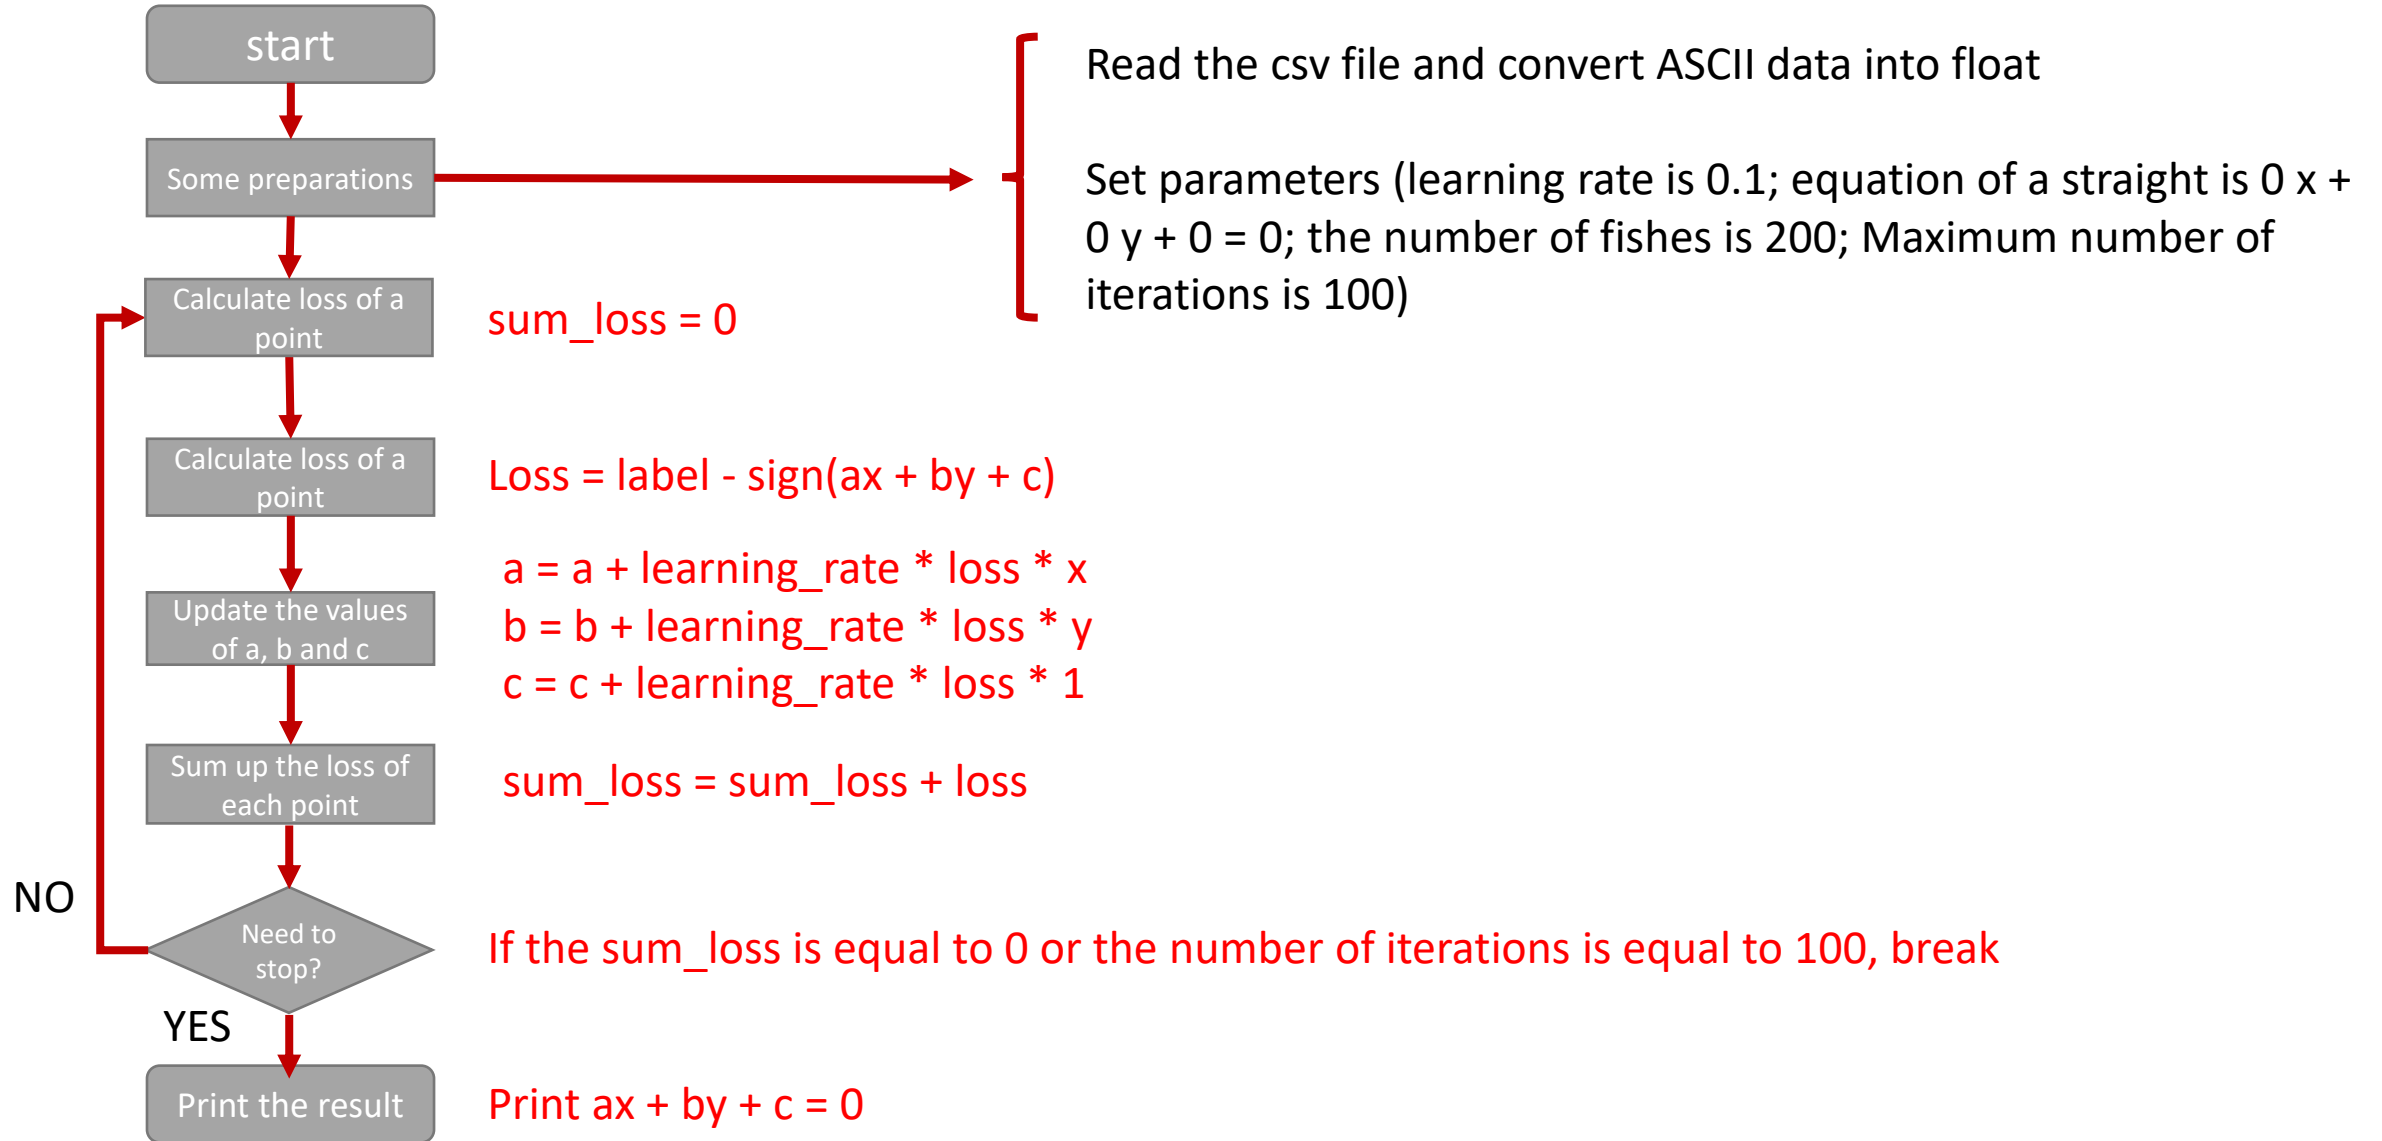
Any ideas about AI?

Assignment

Use perceptron to draw a straight line which can classify salmon and seabass based on length and color: I will provide you with a csv file which contains the length, color and label (-1 or 1) of 200 fishes. Use C to implement a perceptron and print the final equation of the straight line($a x + b y + c = 0$)

- a) The csv file has been uploaded on bb
- b) If the label of a fish is 1, it means this is a salmon; if the label of a fish is -1, it means this is a seabass;
- c) If you don't know how to read a csv file, you can refer to the answer of the previous assignment
- d) If you want to plot the result in an image, you can install OpenCV.
- e) The initial equation of a straight line is : $0 x + 0 y + 0 = 0$
- f) Use `sign()` as the activation function
- g) The learning rate is 0.1

Assignment



Assignment

You can choose to use this function to plot the result

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <math.h>
#include <opencv2/opencv.hpp>

#define num_of_points 200
#define learning_rate 0.1

typedef struct
{
    float x[2];
    int label;
}Point;

Point points[num_of_points];
```


```
void show_result(Point* points, int num, float a, float b, float c)
{
    cv::Mat result = cv::Mat(750, 1500, CV_8UC3, cv::Scalar(0, 0, 0));
    for (int i = 0; i < num; i++)
    {
        cv::Point2f fish_point = cv::Point2f(points[i].x[0] * 50,
        points[i].x[1] * 50);
        if (points[i].label == 1)
            cv::circle(result, fish_point, 5, cv::Scalar(255, 0, 0), -1);
        else
            cv::circle(result, fish_point, 5, cv::Scalar(0, 0, 255), -1);
    }

    cv::line(result, cv::Point2f(0, -1 * c * 50 / b), cv::Point2f(1500, -1
    * (a * 1500 + c * 50) / b), cv::Scalar(0, 255, 0), 2);
    std::cout << (float)c / b << " " << (-1 * a * 1400 + c) / b;
    imshow("result", result);

    cv::waitKey(0);
}
```

Assignment

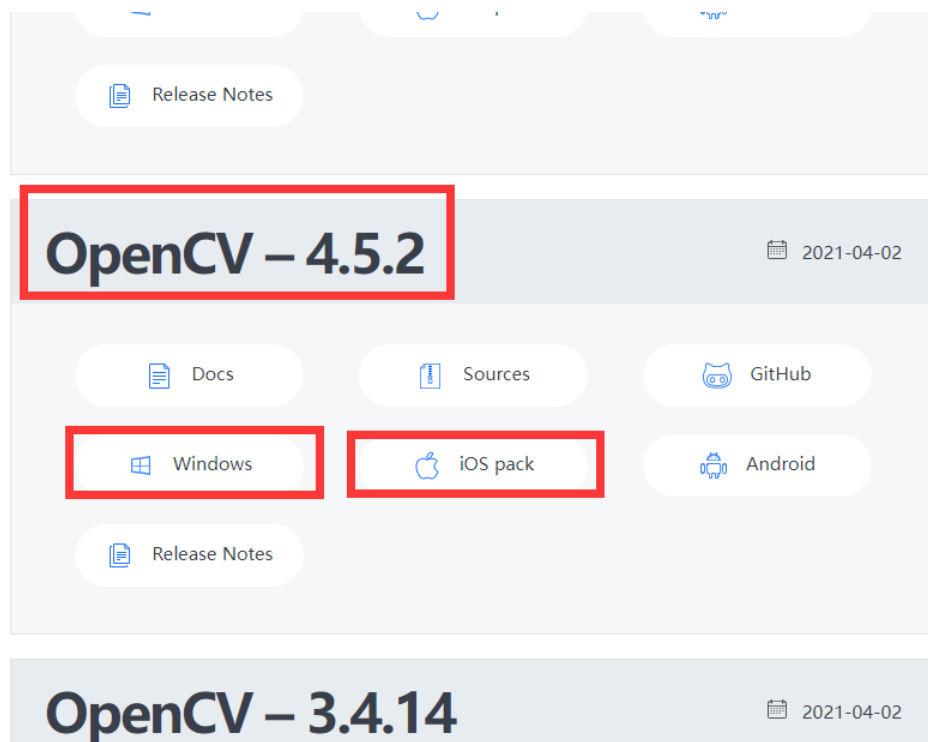
The result looks like this



Install OpenCV on VS 2022

点击这个链接，到opencv官网下载

[Releases - OpenCV](#)



推荐4.5.2版本

根据自己的系统选择安装包

Install OpenCV on VS 2022


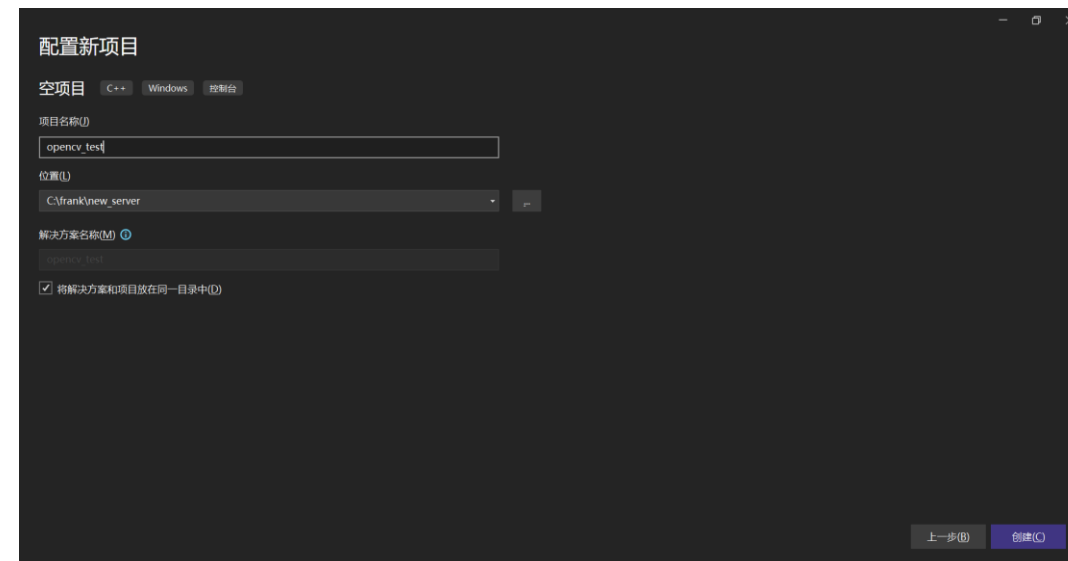
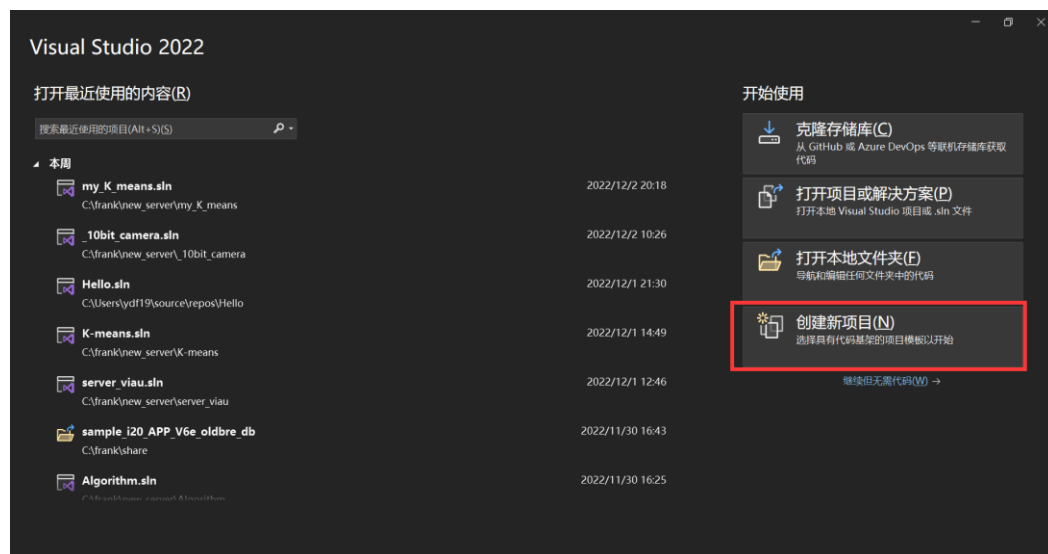


下载完是一个exe文件

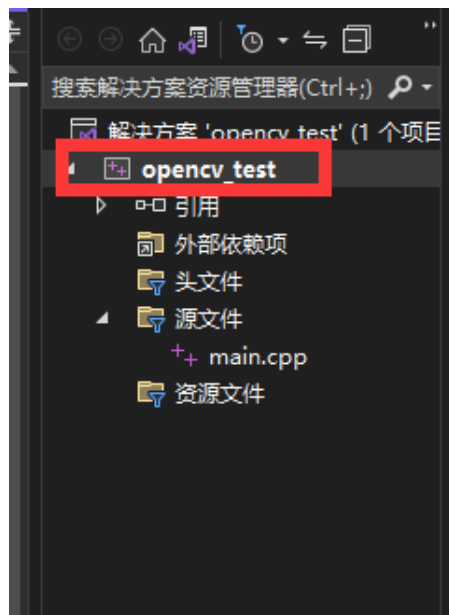
双击后选择一个不包含中文的路径

然后像以前一样新建一个工程

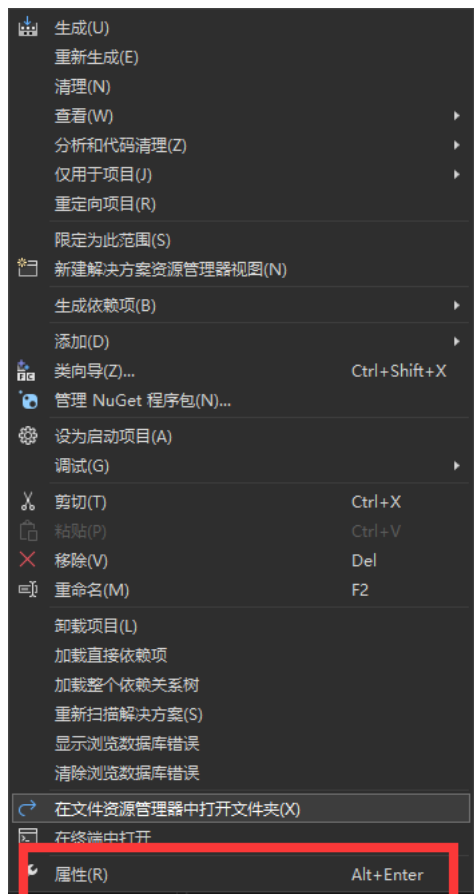
Install OpenCV on VS 2022



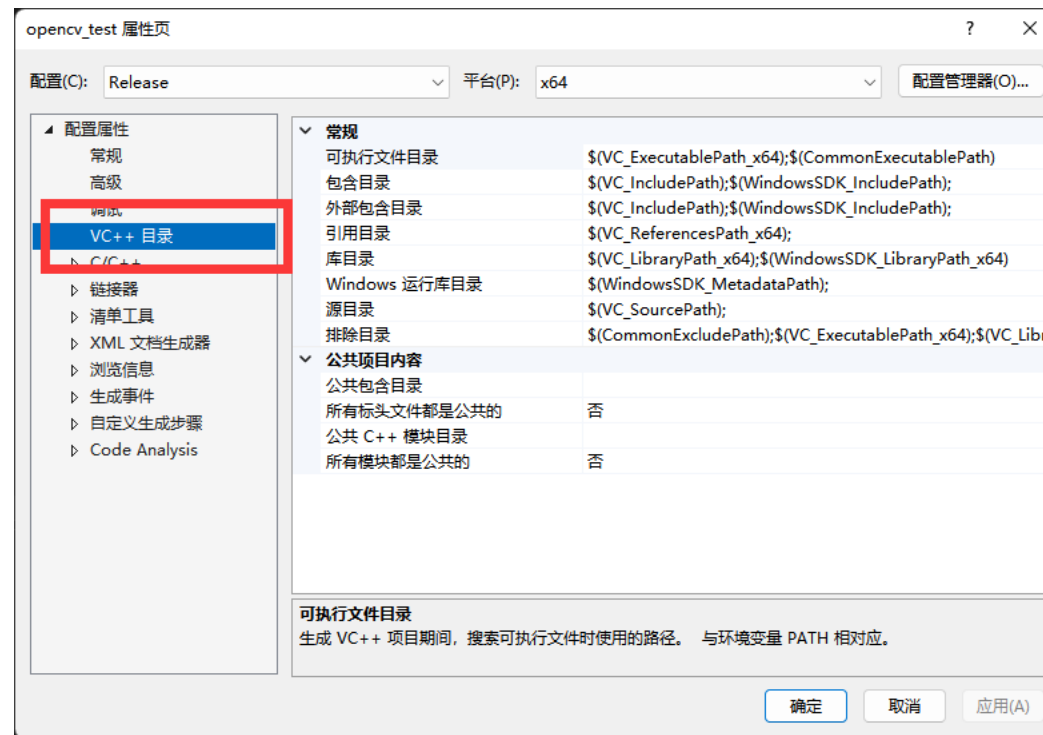
Install OpenCV on VS 2022



右键新建的项目

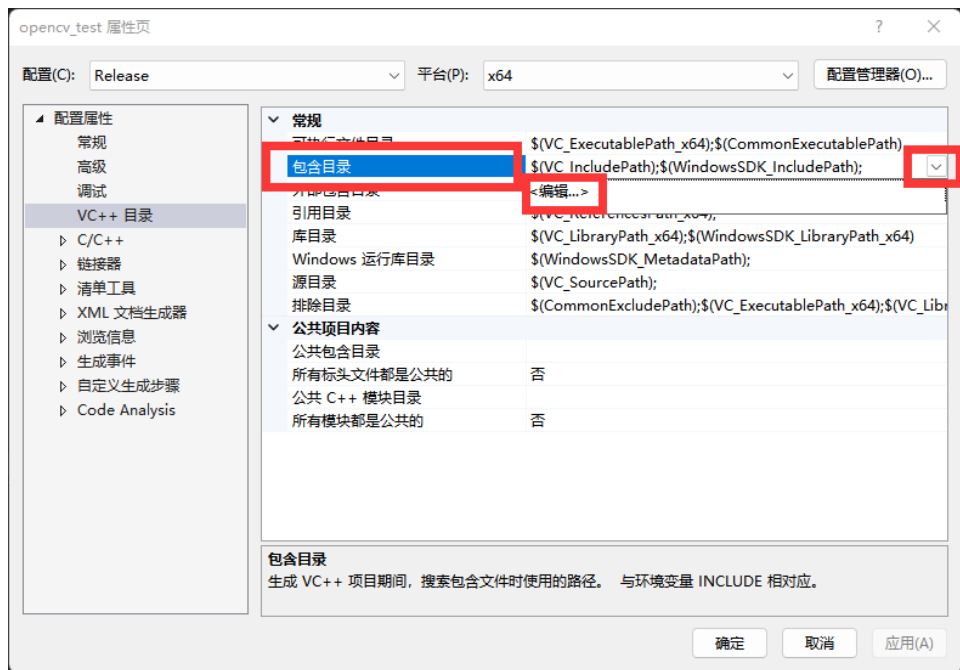


点击属性

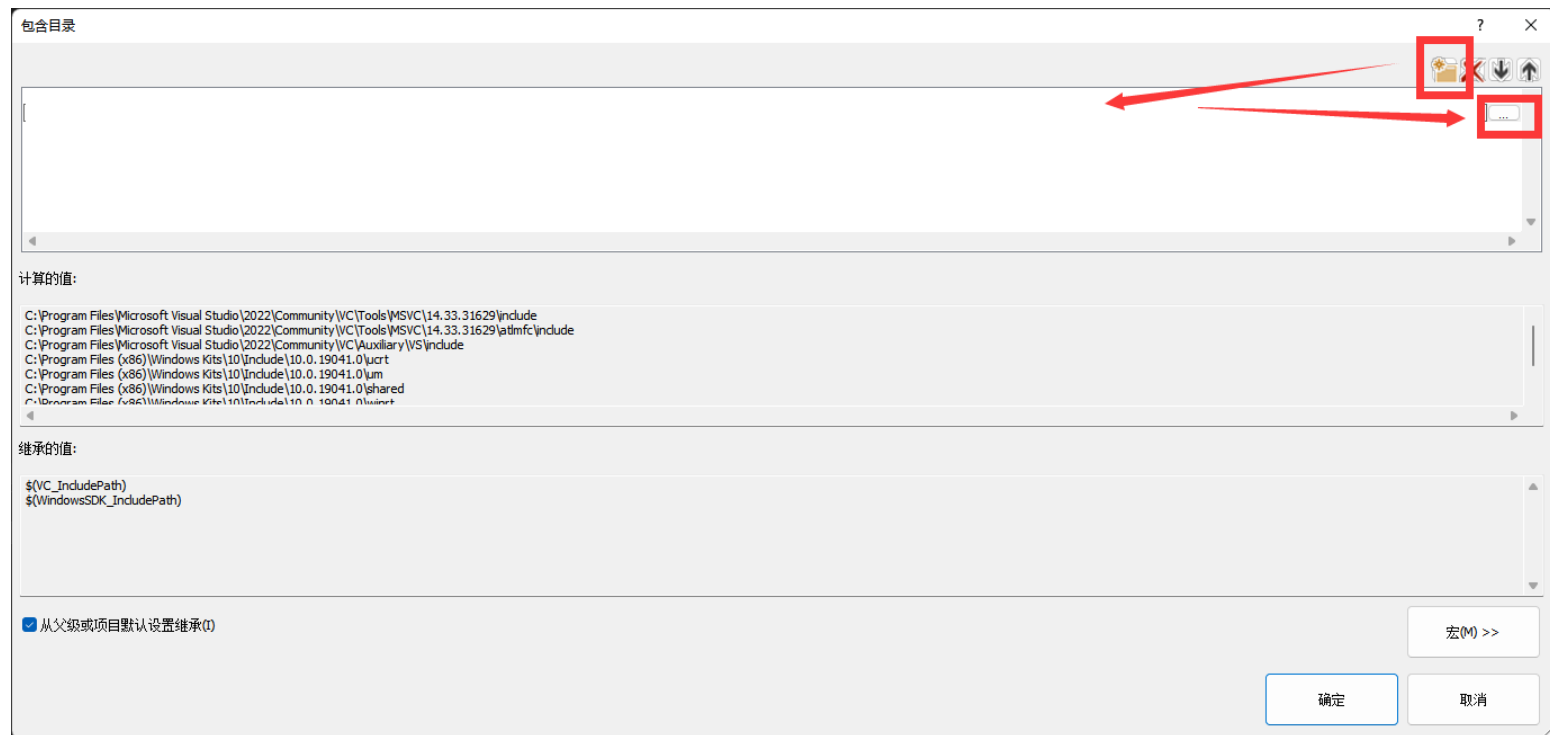


点击VC++ 目录

Install OpenCV on VS 2022

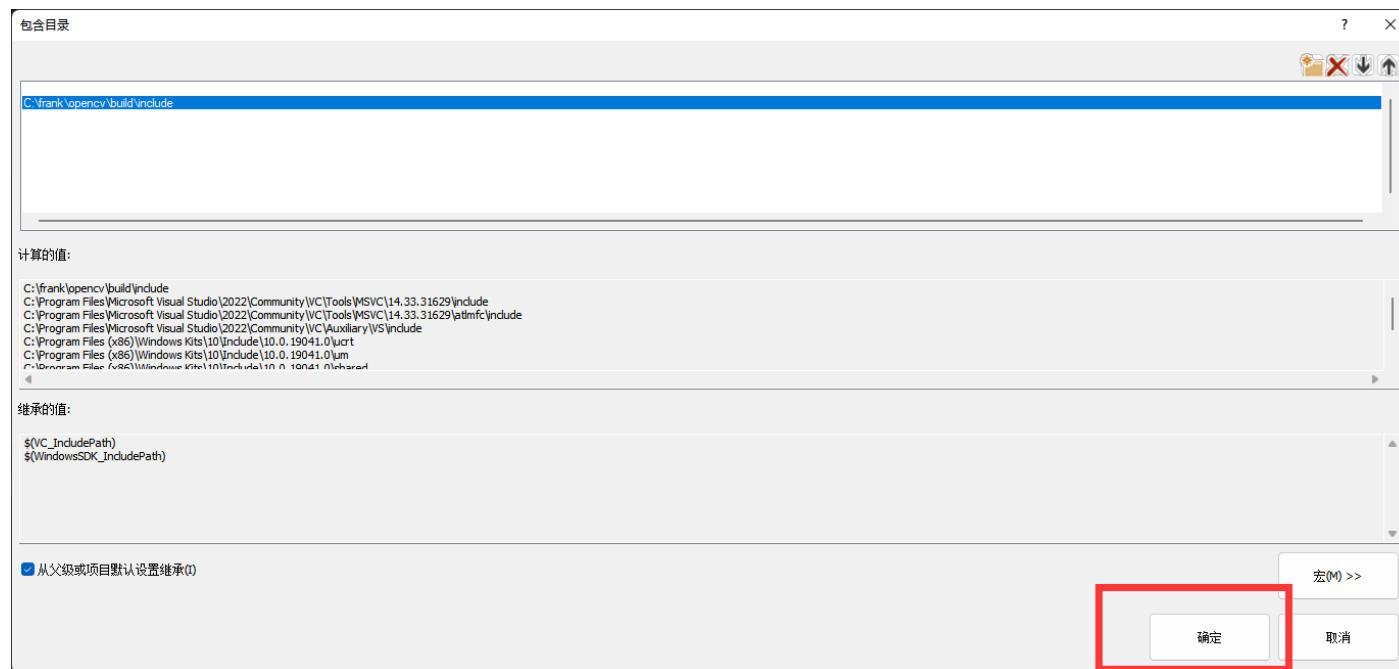
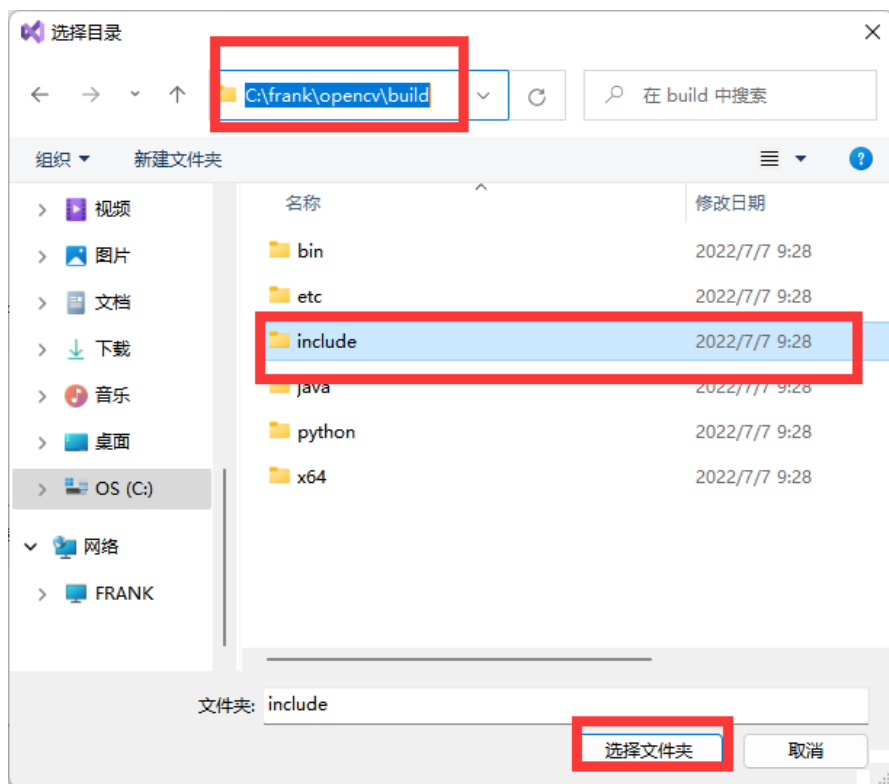


[点击包含目录](#)
[点击右边的↓](#)
[点击编辑](#)



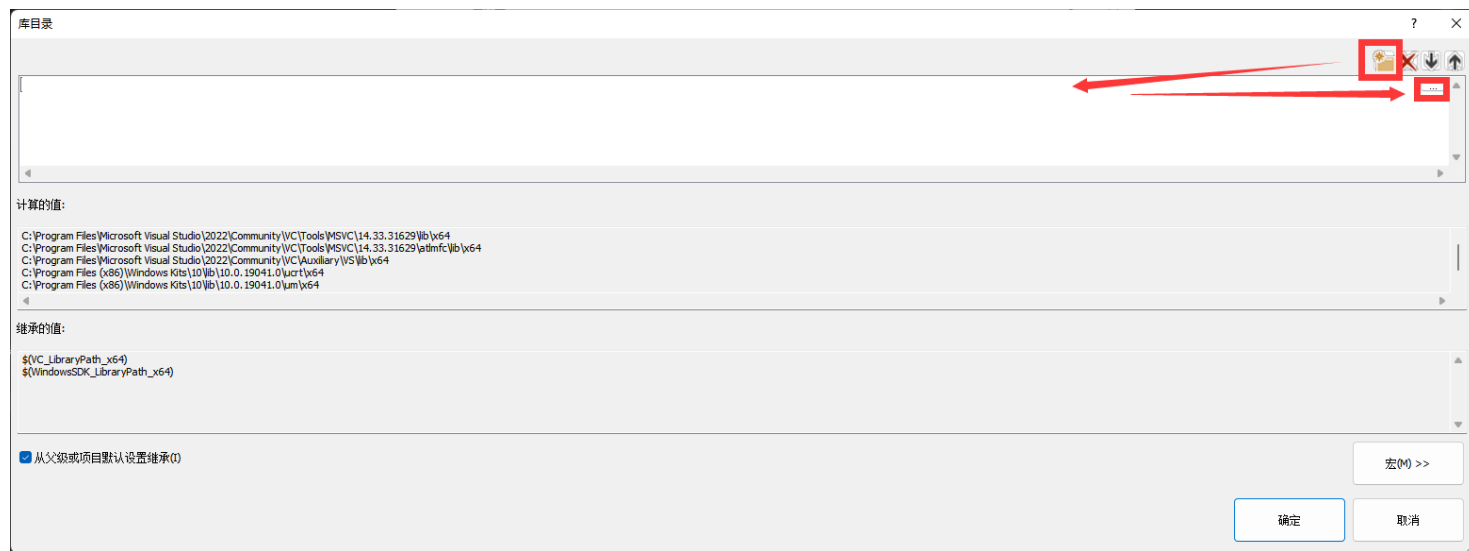
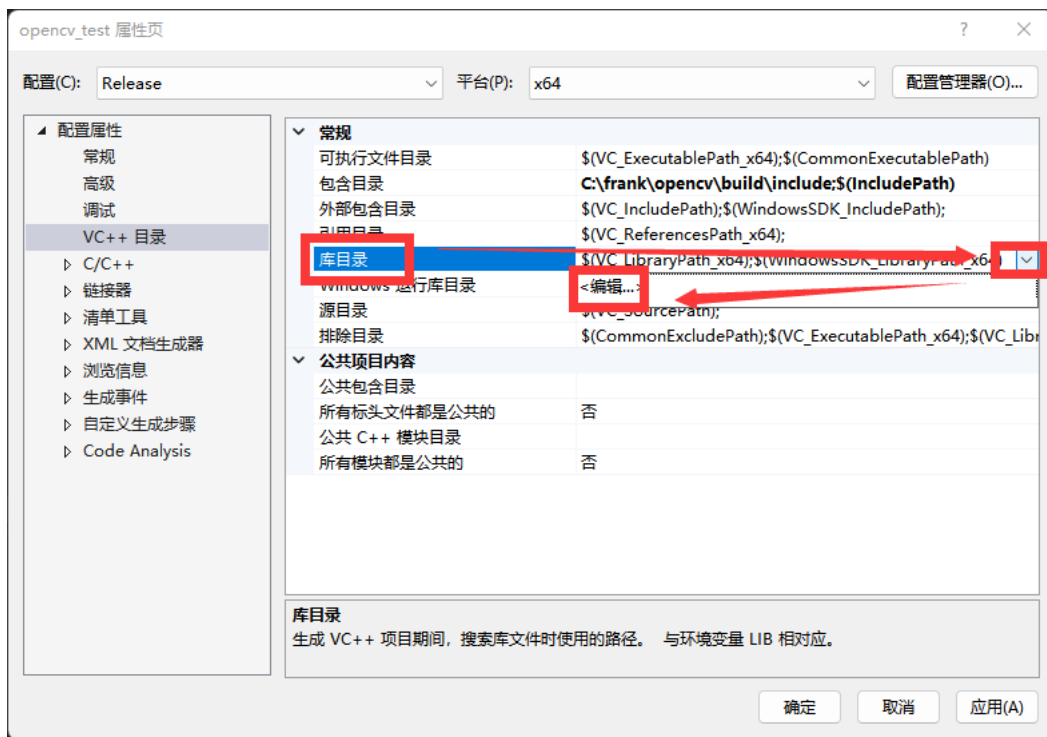
依次点击

Install OpenCV on VS 2022



选择你安装opencv的路径 opencv/build/include 然后确定

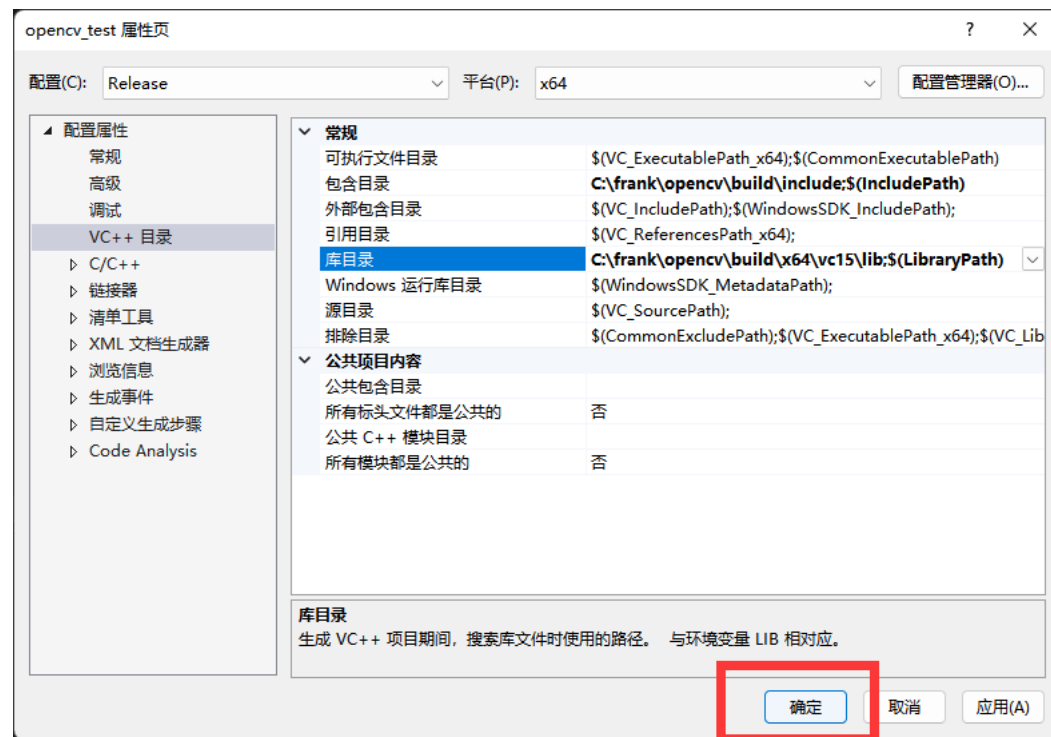
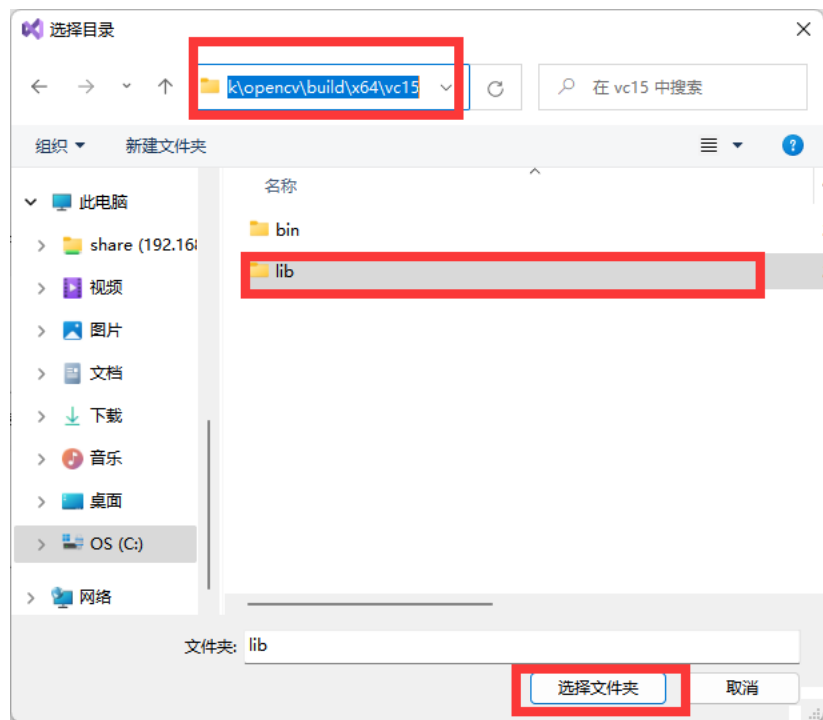
Install OpenCV on VS 2022



依次点击

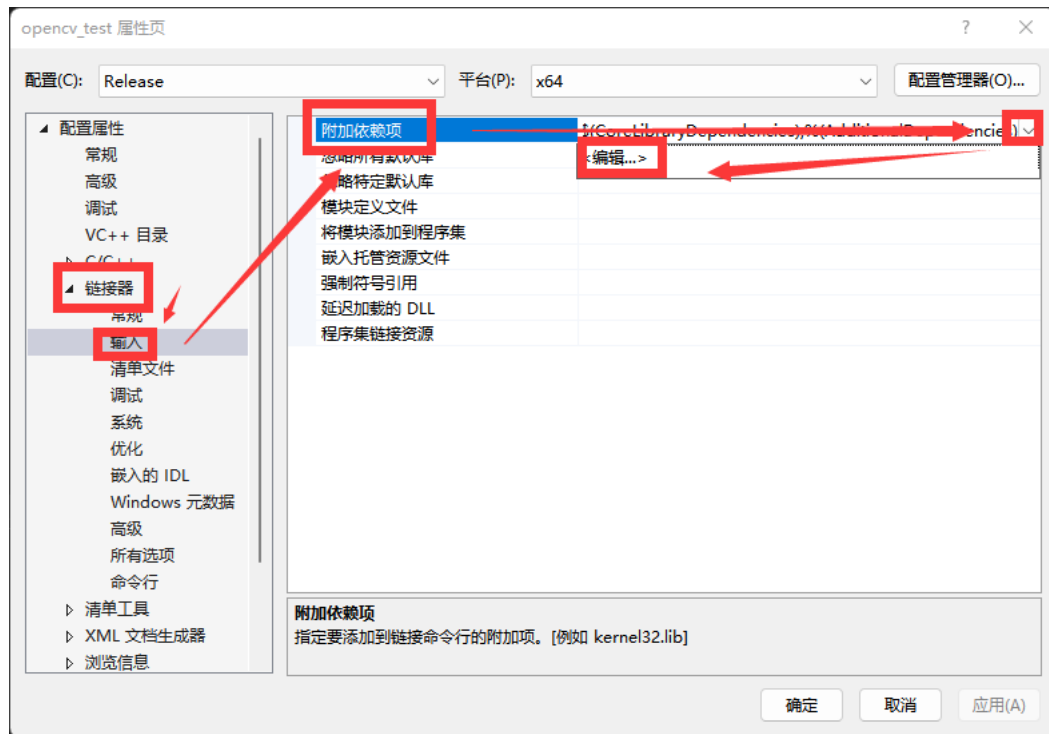
点击库目录 点击右边的↓ 点击编辑

Install OpenCV on VS 2022

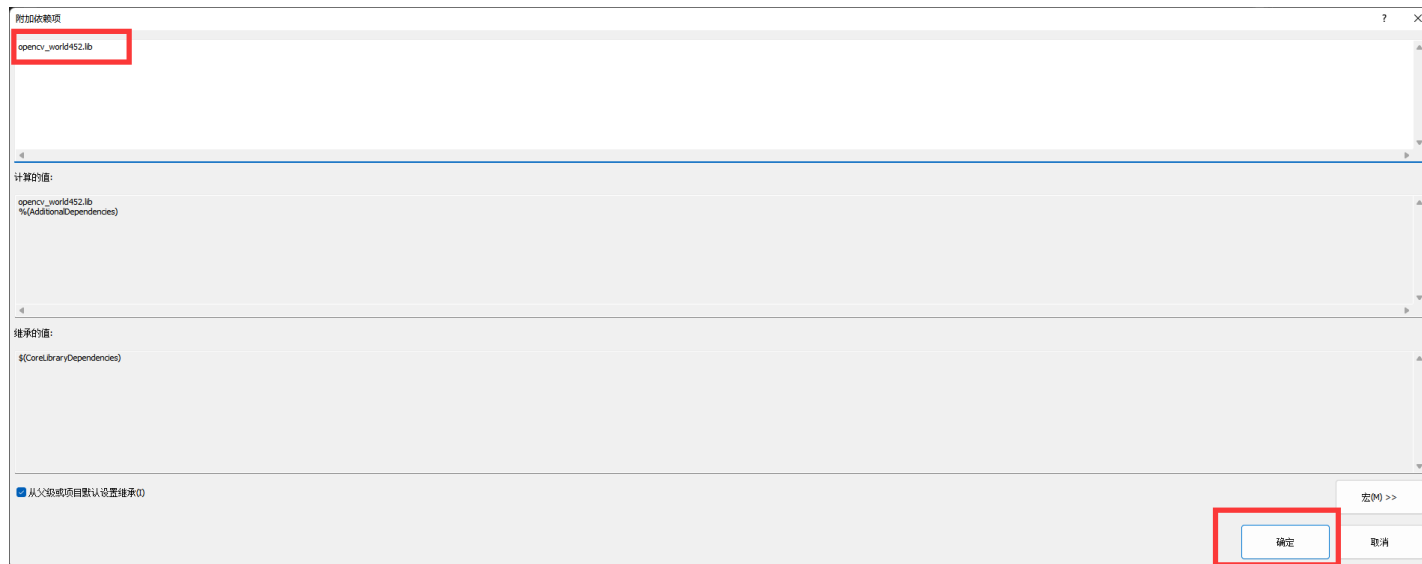


选择你安装opencv的路径 opencv/build/x64/vc15/lib 然后确定

Install OpenCV on VS 2022




选择 链接器 -> 输入 -> 附加依赖项 -> ↓ -> 编辑



手动输入 opencv_world452.lib 然后确定

Install OpenCV on VS 2022



这里改成release x64

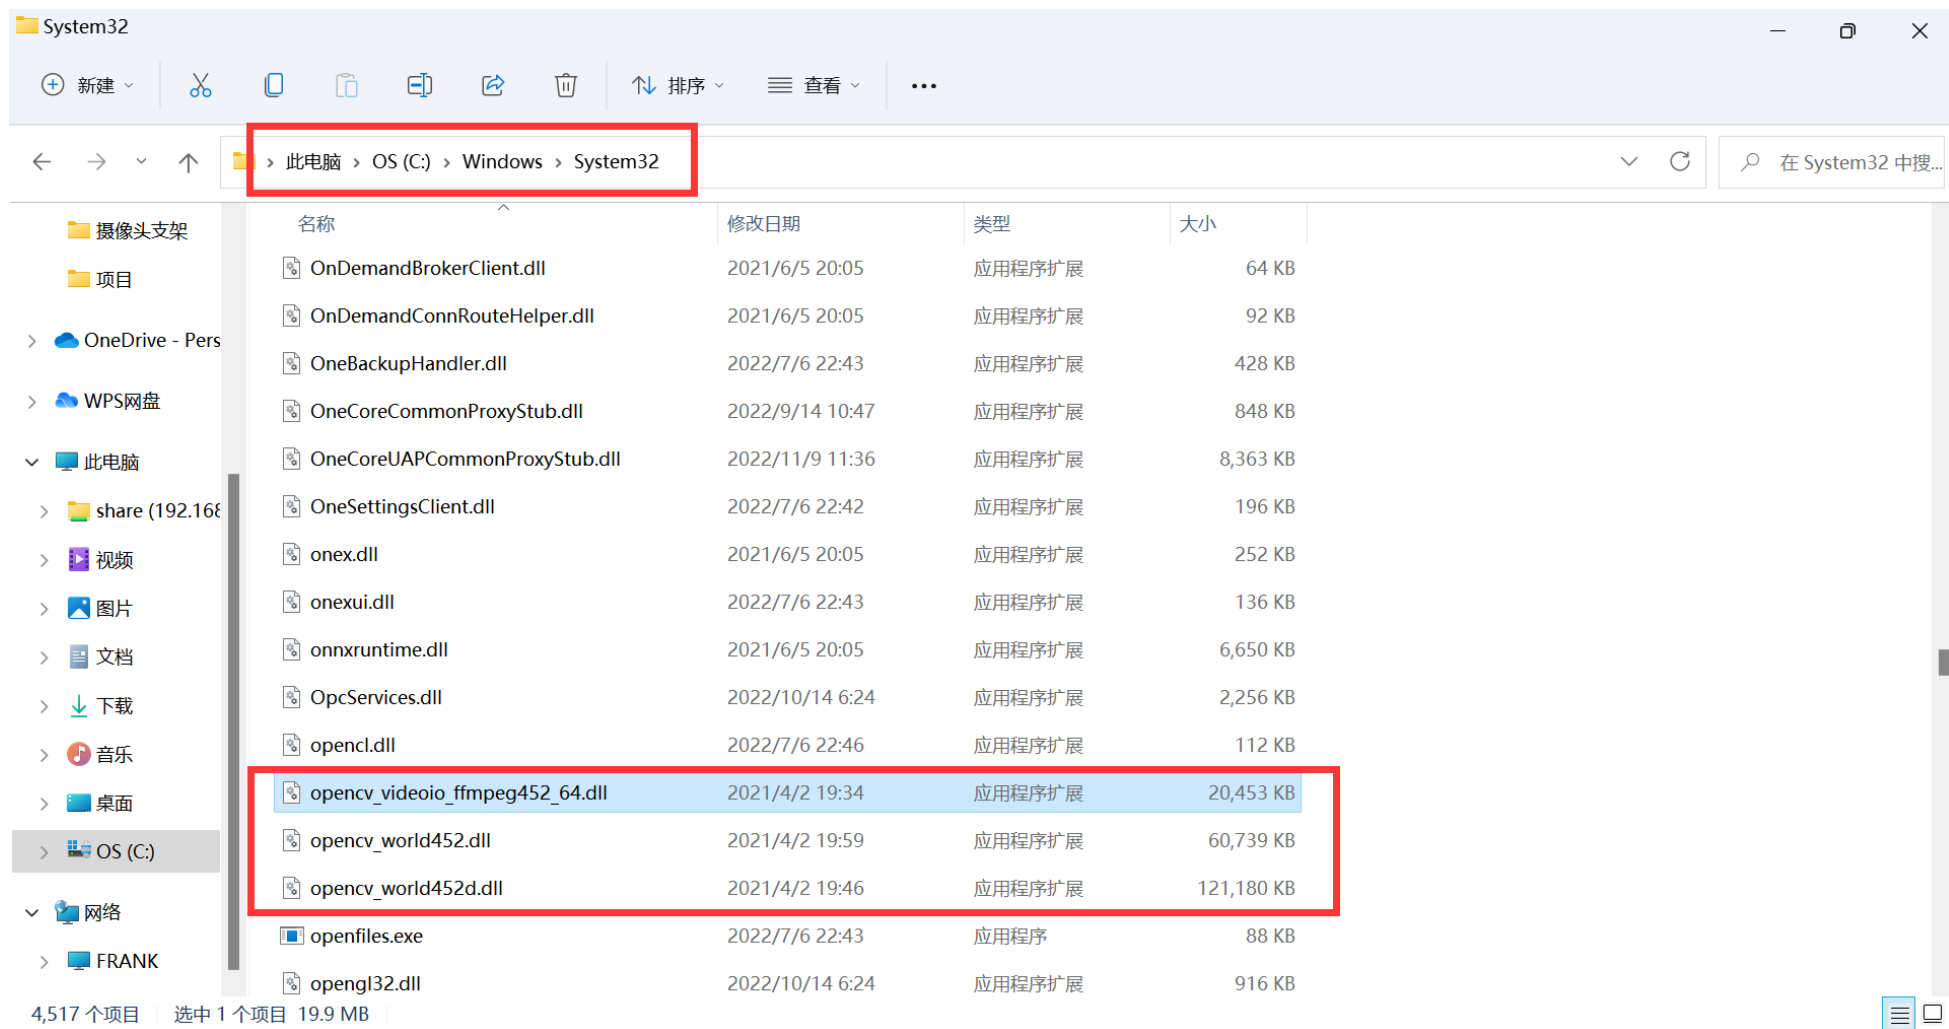
↑ 此电脑 > OS (C:) > frank > opencv > build > x64 > vc15 > bin

名称	修改日期	类型	大小
opencv_annotation.exe	2021/4/2 20:00	应用程序	45 KB
opencv_interactive-calibration.exe	2021/4/2 20:00	应用程序	121 KB
opencv_model_diagnostics.exe	2021/4/2 20:00	应用程序	20 KB
opencv_version.exe	2021/4/2 20:00	应用程序	36 KB
opencv_version_win32.exe	2021/4/2 20:00	应用程序	34 KB
opencv_videoio_ffmpeg452_64.dll	2021/4/2 19:35	应用程序扩展	20,453 KB
opencv_videoio_msmf452_64.dll	2021/4/2 20:00	应用程序扩展	140 KB
opencv_videoio_msmf452_64d.dll	2021/4/2 19:45	应用程序扩展	547 KB
opencv_visualisation.exe	2021/4/2 20:00	应用程序	58 KB
opencv_world452.dll	2021/4/2 20:00	应用程序扩展	60,166 KB
opencv_world452.pdb	2021/4/2 20:00	VisualStudio.pdb.7...	33,852 KB
opencv_world452d.dll	2021/4/2 19:45	应用程序扩展	118,904 KB
opencv_world452d.pdb	2021/4/2 19:45	VisualStudio.pdb.7...	230,436 KB

进入到opencv/build/x64/vc15/bin

把这三个 .dll 文件复制到C: //Windows/system32下

Install OpenCV on VS 2022



就像这样

Install OpenCV on VS 2022

```
#include <iostream>
#include <opencv2/opencv.hpp>
using namespace std;
using namespace cv;
int main()
{
    // 读取图片（使用图片的绝对路径，参考自己的图所在目录）
    Mat srcImg = imread("C://Users//ydf19//Desktop//lenna.jpg");
    if (srcImg.empty()) {
        cout << "could not load image..." << endl;
        return -1;
    }

    imshow("Test opencv setup", srcImg);
    // 显示灰度图
    Mat Gray;
    cvtColor(srcImg, Gray, 6);
    imshow("Gray", Gray);
    // 等待任意按键按下，不添加此语句图片会一闪而过
    waitKey(0);


    cout << "Hello, world." << endl;
    return 0;
}
```

写几行代码测试一下

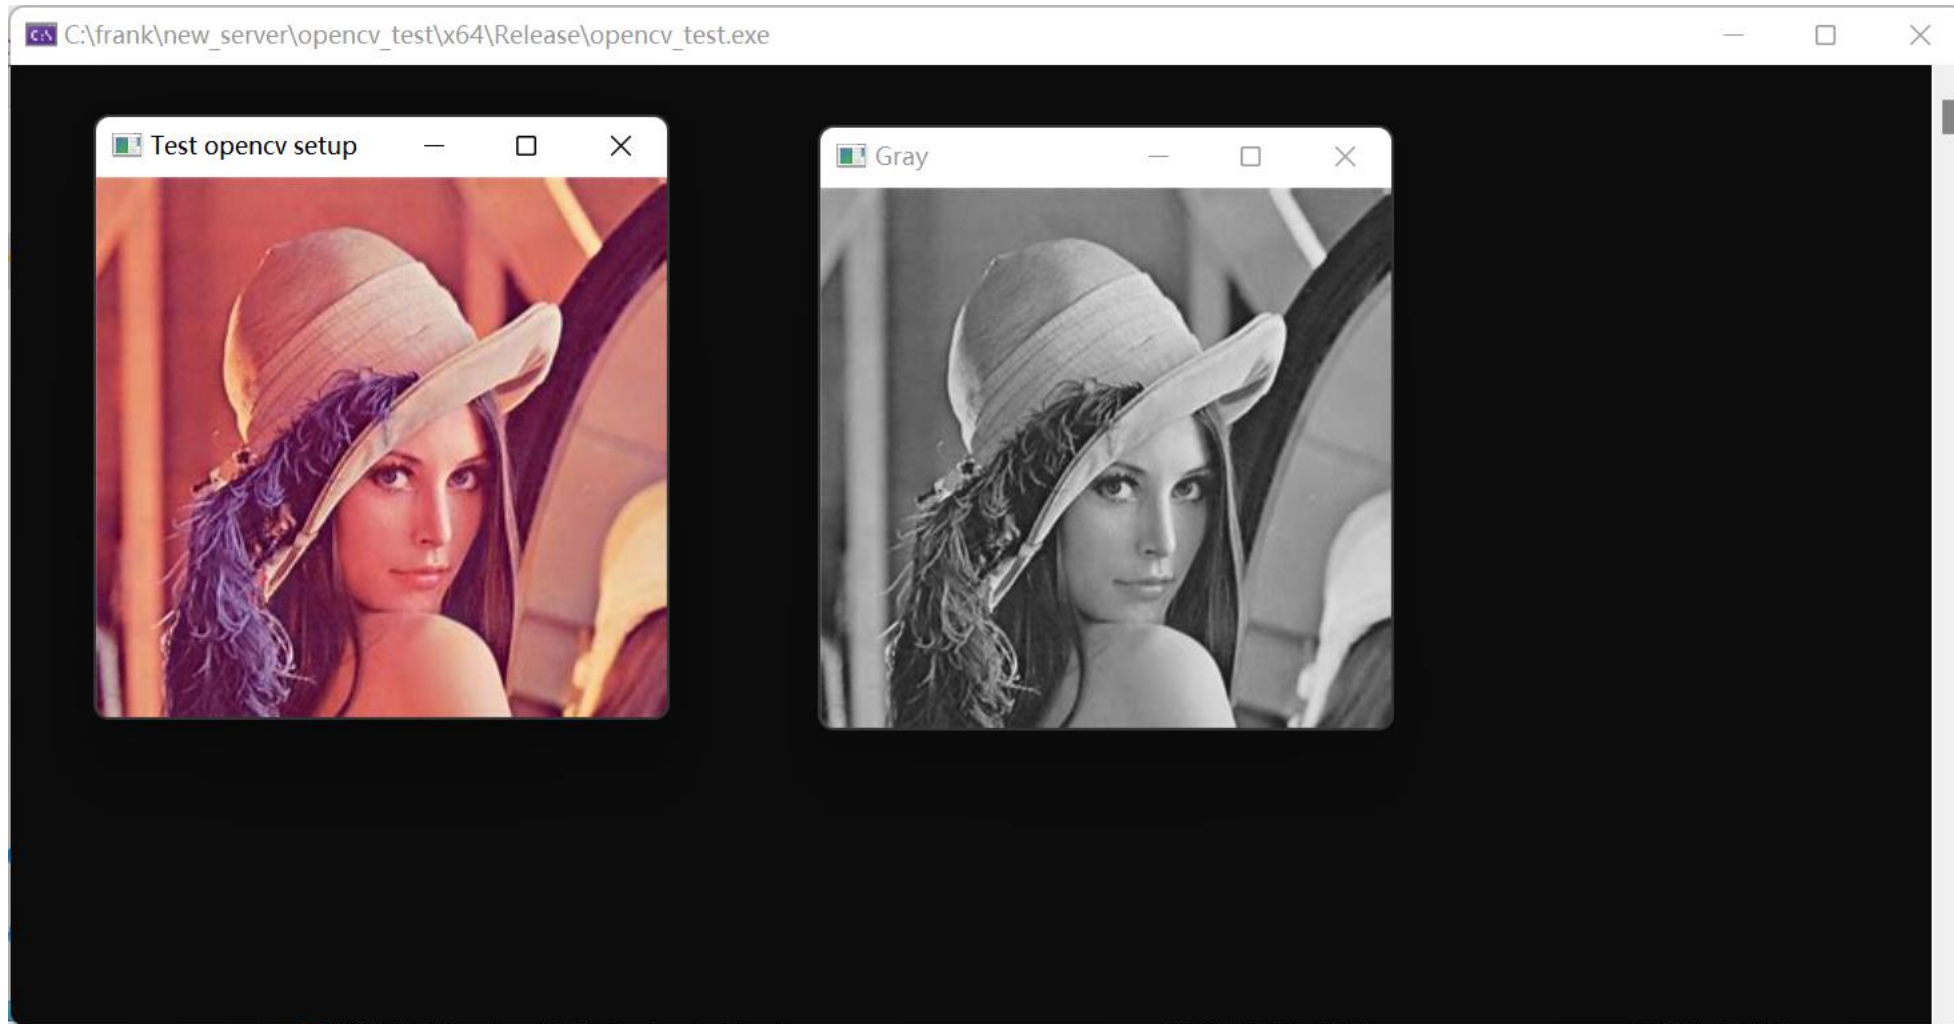
在桌面放一张图片

然后粘贴这些代码

运行



Install OpenCV on VS 2022



成功了

其他系统和编译器的安装可以参考以下链接

Vscode & macOS

https://blog.csdn.net/weixin_43562948/article/details/103956901

Vscode & windows

<https://blog.csdn.net/Avrilzyx/article/details/107036375>

DEV c++ & windows

<https://blog.csdn.net/wadefelix/article/details/1334515>